

Mel-Spectrogram-Only CNN

March 21, 2025

1 CNN for Mel-Spectrogram (30 secs)

1.1 1 - All the imports

```
[1]: import os
import numpy as np
from sklearn.model_selection import train_test_split
import tensorflow as tf
```

2025-03-21 00:51:40.647571: I tensorflow/core/platform/cpu_feature_guard.cc:210] This TensorFlow binary is optimized to use available CPU instructions in performance-critical operations.

To enable the following instructions: AVX2 FMA, in other operations, rebuild TensorFlow with the appropriate compiler flags.

1.2 2 - Put the data within the model

```
[2]: # Import a single image and save it to be read by the model

image = os.path.join('blues.00000.png')

# Load the image
image = tf.io.read_file(image)

# Convert to a numpy array
image = tf.image.decode_png(image, channels=1)
image = tf.image.convert_image_dtype(image, tf.float32)
image = tf.image.resize(image, [256, 256])
image = image.numpy()
```

2 3 - Create the model

```
[3]: from tensorflow.keras import models
from tensorflow.keras.layers import Conv2D, MaxPooling2D, Flatten, Dense, \
↳ Dropout, Normalization

model = models.Sequential([
    Conv2D(32, (3, 3), activation='relu', input_shape=(256, 256, 1)),
```

```

Normalization(),
MaxPooling2D((2, 2)),

Conv2D(64, (3, 3), activation='relu'),
Normalization(),
MaxPooling2D((2, 2)),

Conv2D(128, (3, 3), activation='relu'),
Normalization(),
MaxPooling2D((2, 2)),

Conv2D(256, (3, 3), activation='relu'),
Normalization(),
MaxPooling2D((2, 2)),

Flatten(),

Dense(512, activation='relu'),
Dropout(0.5),

Dense(256, activation='relu'),
Dropout(0.5),

Dense(128, activation='relu'),
Dense(10, activation='softmax')
])

```

/opt/conda/lib/python3.12/site-packages/keras/src/layers/convolutional/base_conv.py:107: UserWarning: Do not pass an `input_shape`/`input_dim` argument to a layer. When using Sequential models, prefer using an `Input(shape)` object as the first layer in the model instead.

```
super().__init__(activity_regularizer=activity_regularizer, **kwargs)
```

```
[4]: model.summary()
```

Model: "sequential"

| Layer (type) | Output Shape | Param # |
|-------------------------------|----------------------|---------|
| conv2d (Conv2D) | (None, 254, 254, 32) | 320 |
| normalization (Normalization) | (None, 254, 254, 32) | 65 |
| max_pooling2d (MaxPooling2D) | (None, 127, 127, 32) | 0 |
| conv2d_1 (Conv2D) | (None, 125, 125, 64) | 18,496 |

| | | |
|---------------------------------|----------------------|------------|
| normalization_1 (Normalization) | (None, 125, 125, 64) | 129 |
| max_pooling2d_1 (MaxPooling2D) | (None, 62, 62, 64) | 0 |
| conv2d_2 (Conv2D) | (None, 60, 60, 128) | 73,856 |
| normalization_2 (Normalization) | (None, 60, 60, 128) | 257 |
| max_pooling2d_2 (MaxPooling2D) | (None, 30, 30, 128) | 0 |
| conv2d_3 (Conv2D) | (None, 28, 28, 256) | 295,168 |
| normalization_3 (Normalization) | (None, 28, 28, 256) | 513 |
| max_pooling2d_3 (MaxPooling2D) | (None, 14, 14, 256) | 0 |
| flatten (Flatten) | (None, 50176) | 0 |
| dense (Dense) | (None, 512) | 25,690,624 |
| dropout (Dropout) | (None, 512) | 0 |
| dense_1 (Dense) | (None, 256) | 131,328 |
| dropout_1 (Dropout) | (None, 256) | 0 |
| dense_2 (Dense) | (None, 128) | 32,896 |
| dense_3 (Dense) | (None, 10) | 1,290 |

Total params: 26,244,942 (100.12 MB)

Trainable params: 26,243,978 (100.11 MB)

Non-trainable params: 964 (3.78 KB)

3 4 - Load the images

```
[5]: import os
import numpy as np
import tensorflow as tf
from sklearn.model_selection import train_test_split
```

```

# Augmentation function
def augment_image(image):
    image = tf.image.random_flip_left_right(image)
    image = tf.image.random_brightness(image, max_delta=0.1)
    image = tf.image.random_contrast(image, 0.8, 1.2)
    return image

# Define the genres and file paths
GENRES = ['blues', 'classical', 'country', 'disco', 'hiphop', 'jazz', 'metal', 'pop', 'reggae', 'rock']
FILE_PATH = os.path.join('Data', 'mel_spectrograms', 'mel_spectrogram_32')
X = []
y = []

GENRE_TO_INDEX = {genre: index for index, genre in enumerate(GENRES)}

# Loop through the genres and load the images with augmentation
for genre in GENRES:
    genre_dir = os.path.join(FILE_PATH, genre)
    print(f"Going through {genre}")
    for file in os.listdir(genre_dir):
        image = tf.io.read_file(os.path.join(genre_dir, file))
        image = tf.image.decode_png(image, channels=1)
        image = tf.image.convert_image_dtype(image, tf.float32)
        image = tf.image.resize(image, [256, 256]) # Resize to 256x256
        image = augment_image(image) # Apply augmentation
        image = image.numpy() # Convert to numpy array
        X.append(image)
        y.append(GENRE_TO_INDEX[genre])

# Convert lists to numpy arrays
X = np.array(X)
y = np.array(y)

# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

```


```

Going through blues
Going through classical
Going through country
Going through disco
Going through hiphop
Going through jazz
Going through metal
Going through pop
Going through reggae

```

Going through rock

```
[6]: # Show image as a sanity check
import matplotlib.pyplot as plt
plt.imshow(X_train[22].reshape(256, 256))
plt.show()
```



3.1 5 - Compile the model

```
[7]: from tensorflow.keras.optimizers import Adam
from tensorflow.keras.callbacks import ReduceLROnPlateau

model.compile(optimizer=Adam(learning_rate=0.0001),
              ↪loss='sparse_categorical_crossentropy', metrics=['accuracy'])
reduce_lr = ReduceLROnPlateau(monitor='val_loss', factor=0.5, patience=3,
                              ↪min_lr=1e-6)
```

3.2 6 - Fit the model

```
[8]: model.fit(X_train, y_train, epochs=20, validation_data=(X_test, y_test),  
             ↪batch_size=32, callbacks=[reduce_lr])
```

```
Epoch 1/20  
25/25          65s 2s/step -  
accuracy: 0.0953 - loss: 2.3052 - val_accuracy: 0.0650 - val_loss: 2.3025 -  
learning_rate: 1.0000e-04  
Epoch 2/20  
25/25          50s 2s/step -  
accuracy: 0.1109 - loss: 2.2979 - val_accuracy: 0.2150 - val_loss: 2.2986 -  
learning_rate: 1.0000e-04  
Epoch 3/20  
25/25          37s 1s/step -  
accuracy: 0.1290 - loss: 2.2930 - val_accuracy: 0.1750 - val_loss: 2.2816 -  
learning_rate: 1.0000e-04  
Epoch 4/20  
25/25          52s 2s/step -  
accuracy: 0.1985 - loss: 2.2460 - val_accuracy: 0.2150 - val_loss: 2.1861 -  
learning_rate: 1.0000e-04  
Epoch 5/20  
25/25          82s 2s/step -  
accuracy: 0.2158 - loss: 2.1599 - val_accuracy: 0.2950 - val_loss: 2.0833 -  
learning_rate: 1.0000e-04  
Epoch 6/20  
25/25          60s 2s/step -  
accuracy: 0.2469 - loss: 2.0954 - val_accuracy: 0.2800 - val_loss: 1.9965 -  
learning_rate: 1.0000e-04  
Epoch 7/20  
25/25         101s 3s/step -  
accuracy: 0.3163 - loss: 1.9898 - val_accuracy: 0.2850 - val_loss: 1.9485 -  
learning_rate: 1.0000e-04  
Epoch 8/20  
25/25          76s 3s/step -  
accuracy: 0.3015 - loss: 1.8769 - val_accuracy: 0.2900 - val_loss: 1.8989 -  
learning_rate: 1.0000e-04  
Epoch 9/20  
25/25          77s 3s/step -  
accuracy: 0.3454 - loss: 1.7936 - val_accuracy: 0.3000 - val_loss: 1.8894 -  
learning_rate: 1.0000e-04  
Epoch 10/20  
25/25          74s 3s/step -  
accuracy: 0.3017 - loss: 1.8817 - val_accuracy: 0.3050 - val_loss: 1.8379 -  
learning_rate: 1.0000e-04  
Epoch 11/20  
25/25          84s 3s/step -  
accuracy: 0.3426 - loss: 1.7582 - val_accuracy: 0.3400 - val_loss: 1.8200 -
```

```

learning_rate: 1.0000e-04
Epoch 12/20
25/25          101s 4s/step -
accuracy: 0.3505 - loss: 1.7707 - val_accuracy: 0.3100 - val_loss: 1.8360 -
learning_rate: 1.0000e-04
Epoch 13/20
25/25          85s 3s/step -
accuracy: 0.3476 - loss: 1.7588 - val_accuracy: 0.3150 - val_loss: 1.8088 -
learning_rate: 1.0000e-04
Epoch 14/20
25/25          147s 4s/step -
accuracy: 0.3952 - loss: 1.6645 - val_accuracy: 0.3650 - val_loss: 1.7219 -
learning_rate: 1.0000e-04
Epoch 15/20
25/25          67s 3s/step -
accuracy: 0.3605 - loss: 1.7098 - val_accuracy: 0.3550 - val_loss: 1.7054 -
learning_rate: 1.0000e-04
Epoch 16/20
25/25          66s 3s/step -
accuracy: 0.3802 - loss: 1.6301 - val_accuracy: 0.3500 - val_loss: 1.7528 -
learning_rate: 1.0000e-04
Epoch 17/20
25/25          76s 3s/step -
accuracy: 0.3605 - loss: 1.7336 - val_accuracy: 0.3850 - val_loss: 1.6834 -
learning_rate: 1.0000e-04
Epoch 18/20
25/25          72s 3s/step -
accuracy: 0.4093 - loss: 1.6555 - val_accuracy: 0.4350 - val_loss: 1.6226 -
learning_rate: 1.0000e-04
Epoch 19/20
25/25          81s 3s/step -
accuracy: 0.4258 - loss: 1.5318 - val_accuracy: 0.4100 - val_loss: 1.6149 -
learning_rate: 1.0000e-04
Epoch 20/20
25/25          73s 3s/step -
accuracy: 0.4750 - loss: 1.4560 - val_accuracy: 0.3900 - val_loss: 1.6673 -
learning_rate: 1.0000e-04

```

[8]: <keras.src.callbacks.history.History at 0x7fcab0128ec0>

3.3 7 - Check the accuracy

```

[9]: evaluation = model.evaluate(X_test, y_test)
      print(f"Test accuracy: {evaluation[1]:.3f}")

```

```

7/7          4s 574ms/step -
accuracy: 0.3926 - loss: 1.6304
Test accuracy: 0.390

```

4 8 - Apply the confusion matrix after the model


```
[10]: import seaborn as sns
# from sklearn.metrics import confusion
import numpy as NP
from sklearn.metrics import confusion_matrix

cnn_preds = np.argmax(model.predict(X_test), axis=1)
cnn_cm = confusion_matrix(y_test, cnn_preds)

# Plot the confusion matrix
plt.figure(figsize=(8, 6))
sns.heatmap(cnn_cm, annot=True, fmt="d", cmap="Blues", xticklabels=GENRES, yticklabels=GENRES)
plt.title("CNN Confusion Matrix")
plt.xlabel("Predicted")
plt.ylabel("True")
plt.show()
```

7/7

5s 658ms/step



4.1 9 - Limited Genres Easy (metal and classical)

```
[11]: import os
import numpy as np
import tensorflow as tf
from sklearn.model_selection import train_test_split

# Augmentation function
def augment_image(image):
    image = tf.image.random_flip_left_right(image)
    image = tf.image.random_brightness(image, max_delta=0.1)
    image = tf.image.random_contrast(image, 0.8, 1.2)
    return image

# Define the genres and file paths
GENRES = ['classical', 'metal']
FILE_PATH = os.path.join('Data', 'mel_spectrograms', 'mel_spectrogram_32')
X = []
y = []

GENRE_TO_INDEX = {genre: index for index, genre in enumerate(GENRES)}

# Loop through the genres and load the images with augmentation
for genre in GENRES:
    genre_dir = os.path.join(FILE_PATH, genre)
    print(f"Going through {genre}")
    for file in os.listdir(genre_dir):
        image = tf.io.read_file(os.path.join(genre_dir, file))
        image = tf.image.decode_png(image, channels=1)
        image = tf.image.convert_image_dtype(image, tf.float32)
        image = tf.image.resize(image, [256, 256]) # Resize to 256x256
        image = augment_image(image) # Apply augmentation
        image = image.numpy() # Convert to numpy array
        X.append(image)
        y.append(GENRE_TO_INDEX[genre])

# Convert lists to numpy arrays
X = np.array(X)
y = np.array(y)

# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
    random_state=42)
```

```

from tensorflow.keras import models
from tensorflow.keras.layers import Conv2D, MaxPooling2D, Flatten, Dense,
    ↳Dropout, Normalization

model = models.Sequential([
    Conv2D(32, (3, 3), activation='relu', input_shape=(256, 256, 1)),
    Normalization(),
    MaxPooling2D((2, 2)),

    Conv2D(64, (3, 3), activation='relu'),
    Normalization(),
    MaxPooling2D((2, 2)),

    Conv2D(128, (3, 3), activation='relu'),
    Normalization(),
    MaxPooling2D((2, 2)),

    Conv2D(256, (3, 3), activation='relu'),
    Normalization(),
    MaxPooling2D((2, 2)),

    Flatten(),

    Dense(512, activation='relu'),
    Dropout(0.5),

    Dense(256, activation='relu'),
    Dropout(0.5),

    Dense(128, activation='relu'),
    Dense(10, activation='softmax')
])

from tensorflow.keras.optimizers import Adam
from tensorflow.keras.callbacks import ReduceLROnPlateau

model.compile(optimizer=Adam(learning_rate=0.0001),
    ↳loss='sparse_categorical_crossentropy', metrics=['accuracy'])
reduce_lr = ReduceLROnPlateau(monitor='val_loss', factor=0.5, patience=3,
    ↳min_lr=1e-6)

model.fit(X_train, y_train, epochs=20, validation_data=(X_test, y_test),
    ↳batch_size=32, callbacks=[reduce_lr])

evaluation = model.evaluate(X_test, y_test)
print(f"Test accuracy: {evaluation[1]:.3f}")

```

Going through classical

Going through metal

/opt/conda/lib/python3.12/site-

packages/keras/src/layers/convolutional/base_conv.py:107: UserWarning: Do not pass an `input_shape`/`input_dim` argument to a layer. When using Sequential models, prefer using an `Input(shape)` object as the first layer in the model instead.

super().__init__(activity_regularizer=activity_regularizer, **kwargs)

Epoch 1/20

5/5 23s 3s/step -

accuracy: 0.1472 - loss: 2.2249 - val_accuracy: 0.5250 - val_loss: 1.7616 -
learning_rate: 1.0000e-04

Epoch 2/20

5/5 20s 3s/step -

accuracy: 0.4185 - loss: 1.6901 - val_accuracy: 0.5250 - val_loss: 0.9549 -
learning_rate: 1.0000e-04

Epoch 3/20

5/5 21s 3s/step -

accuracy: 0.4997 - loss: 1.1220 - val_accuracy: 0.9500 - val_loss: 0.6474 -
learning_rate: 1.0000e-04

Epoch 4/20

5/5 15s 3s/step -

accuracy: 0.5453 - loss: 1.1086 - val_accuracy: 0.5500 - val_loss: 0.6328 -
learning_rate: 1.0000e-04

Epoch 5/20

5/5 20s 3s/step -

accuracy: 0.5001 - loss: 1.0571 - val_accuracy: 0.9500 - val_loss: 0.6300 -
learning_rate: 1.0000e-04

Epoch 6/20

5/5 21s 3s/step -

accuracy: 0.5686 - loss: 0.9307 - val_accuracy: 0.8500 - val_loss: 0.6722 -
learning_rate: 1.0000e-04

Epoch 7/20

5/5 21s 3s/step -

accuracy: 0.6301 - loss: 0.7912 - val_accuracy: 0.9500 - val_loss: 0.6510 -
learning_rate: 1.0000e-04

Epoch 8/20

5/5 14s 3s/step -

accuracy: 0.8193 - loss: 0.4390 - val_accuracy: 0.8250 - val_loss: 0.3592 -
learning_rate: 1.0000e-04

Epoch 15/20

5/5 20s 3s/step -

accuracy: 0.8481 - loss: 0.3629 - val_accuracy: 0.9500 - val_loss: 0.2301 -
learning_rate: 1.0000e-04

Epoch 16/20

5/5 12s 2s/step -

accuracy: 0.8777 - loss: 0.3015 - val_accuracy: 0.7750 - val_loss: 0.4160 -

```

learning_rate: 1.0000e-04
Epoch 17/20
5/5          23s 3s/step -
accuracy: 0.7872 - loss: 0.4212 - val_accuracy: 0.9250 - val_loss: 0.2505 -
learning_rate: 1.0000e-04
Epoch 18/20
5/5          14s 3s/step -
accuracy: 0.8219 - loss: 0.3761 - val_accuracy: 0.8250 - val_loss: 0.2870 -
learning_rate: 1.0000e-04
Epoch 19/20
5/5          14s 3s/step -
accuracy: 0.9102 - loss: 0.2578 - val_accuracy: 0.9500 - val_loss: 0.2083 -
learning_rate: 5.0000e-05
Epoch 20/20
5/5          14s 3s/step -
accuracy: 0.8877 - loss: 0.3558 - val_accuracy: 0.9750 - val_loss: 0.1807 -
learning_rate: 5.0000e-05
2/2          1s 185ms/step -
accuracy: 0.9729 - loss: 0.1903
Test accuracy: 0.975

```

4.2 10 - Confusion Matrix Easy (metal and classical)

```

[12]: import seaborn as sns
      # from sklearn.metrics import confusion
      import numpy as NP
      from sklearn.metrics import confusion_matrix

      cnn_preds = np.argmax(model.predict(X_test), axis=1)
      cnn_cm = confusion_matrix(y_test, cnn_preds)


      # Plot the confusion matrix
      plt.figure(figsize=(8, 6))
      sns.heatmap(cnn_cm, annot=True, fmt="d", cmap="Blues", xticklabels=GENRES,
                  yticklabels=GENRES)
      plt.title("CNN Confusion Matrix")
      plt.xlabel("Predicted")
      plt.ylabel("True")
      plt.show()

```

```

2/2          1s 533ms/step

```



4.3 11 - Limited genres Hard (disco and pop)

```
[13]: import os
import numpy as np
import tensorflow as tf
from sklearn.model_selection import train_test_split

# Augmentation function
def augment_image(image):
    image = tf.image.random_flip_left_right(image)
    image = tf.image.random_brightness(image, max_delta=0.1)
    image = tf.image.random_contrast(image, 0.8, 1.2)
    return image

# Define the genres and file paths
GENRES = ['disco', 'pop']
FILE_PATH = os.path.join('Data', 'mel_spectrograms', 'mel_spectrogram_32')
```

```

X = []
y = []

GENRE_TO_INDEX = {genre: index for index, genre in enumerate(GENRES)}

# Loop through the genres and load the images with augmentation
for genre in GENRES:
    genre_dir = os.path.join(FILE_PATH, genre)
    print(f"Going through {genre}")
    for file in os.listdir(genre_dir):
        image = tf.io.read_file(os.path.join(genre_dir, file))
        image = tf.image.decode_png(image, channels=1)
        image = tf.image.convert_image_dtype(image, tf.float32)
        image = tf.image.resize(image, [256, 256]) # Resize to 256x256
        image = augment_image(image) # Apply augmentation
        image = image.numpy() # Convert to numpy array
        X.append(image)
        y.append(GENRE_TO_INDEX[genre])

# Convert lists to numpy arrays
X = np.array(X)
y = np.array(y)

# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
    ↪random_state=42)

from tensorflow.keras import models
from tensorflow.keras.layers import Conv2D, MaxPooling2D, Flatten, Dense,
    ↪Dropout, Normalization

model = models.Sequential([
    Conv2D(32, (3, 3), activation='relu', input_shape=(256, 256, 1)),
    Normalization(),
    MaxPooling2D((2, 2)),

    Conv2D(64, (3, 3), activation='relu'),
    Normalization(),
    MaxPooling2D((2, 2)),

    Conv2D(128, (3, 3), activation='relu'),
    Normalization(),
    MaxPooling2D((2, 2)),

    Conv2D(256, (3, 3), activation='relu'),
    Normalization(),
    MaxPooling2D((2, 2)),

```

```

        Flatten(),

        Dense(512, activation='relu'),
        Dropout(0.5),

        Dense(256, activation='relu'),
        Dropout(0.5),

        Dense(128, activation='relu'),
        Dense(10, activation='softmax')
    ])

from tensorflow.keras.optimizers import Adam
from tensorflow.keras.callbacks import ReduceLROnPlateau

model.compile(optimizer=Adam(learning_rate=0.0001),
    ↳loss='sparse_categorical_crossentropy', metrics=['accuracy'])
reduce_lr = ReduceLROnPlateau(monitor='val_loss', factor=0.5, patience=3,
    ↳min_lr=1e-6)

model.fit(X_train, y_train, epochs=20, validation_data=(X_test, y_test),
    ↳batch_size=32, callbacks=[reduce_lr])

evaluation = model.evaluate(X_test, y_test)
print(f"Test accuracy: {evaluation[1]:.3f}")

```

Going through disco

Going through pop

/opt/conda/lib/python3.12/site-

packages/keras/src/layers/convolutional/base_conv.py:107: UserWarning: Do not pass an `input_shape`/`input_dim` argument to a layer. When using Sequential models, prefer using an `Input(shape)` object as the first layer in the model instead.

```
super().__init__(activity_regularizer=activity_regularizer, **kwargs)
```

Epoch 1/20

5/5 22s 3s/step -

accuracy: 0.0952 - loss: 2.2424 - val_accuracy: 0.4750 - val_loss: 1.8354 -
learning_rate: 1.0000e-04

Epoch 2/20

5/5 20s 3s/step -

accuracy: 0.3153 - loss: 1.7927 - val_accuracy: 0.5250 - val_loss: 1.0843 -
learning_rate: 1.0000e-04

Epoch 3/20

5/5 21s 3s/step -

accuracy: 0.3611 - loss: 1.3989 - val_accuracy: 0.5250 - val_loss: 0.7397 -

```

learning_rate: 1.0000e-04
Epoch 4/20
5/5          14s 3s/step -
accuracy: 0.4973 - loss: 1.1316 - val_accuracy: 0.5250 - val_loss: 0.7082 -
learning_rate: 1.0000e-04
Epoch 5/20
5/5          14s 3s/step -
accuracy: 0.5055 - loss: 1.0738 - val_accuracy: 0.5250 - val_loss: 0.7242 -
learning_rate: 1.0000e-04
Epoch 6/20
5/5          20s 3s/step -
accuracy: 0.4804 - loss: 1.1975 - val_accuracy: 0.5250 - val_loss: 0.7851 -
learning_rate: 1.0000e-04
Epoch 7/20
5/5          21s 3s/step -
accuracy: 0.4980 - loss: 0.9551 - val_accuracy: 0.5250 - val_loss: 0.7897 -
learning_rate: 1.0000e-04
Epoch 8/20
5/5          15s 3s/step -
accuracy: 0.4862 - loss: 0.9490 - val_accuracy: 0.5250 - val_loss: 0.7683 -
learning_rate: 5.0000e-05
Epoch 9/20
5/5          20s 3s/step -
accuracy: 0.5511 - loss: 0.9116 - val_accuracy: 0.5250 - val_loss: 0.7425 -
learning_rate: 5.0000e-05
Epoch 10/20
5/5          11s 2s/step -
accuracy: 0.5001 - loss: 0.9749 - val_accuracy: 0.5250 - val_loss: 0.7227 -
learning_rate: 5.0000e-05
Epoch 11/20
5/5          24s 3s/step -
accuracy: 0.4311 - loss: 0.9582 - val_accuracy: 0.5250 - val_loss: 0.7189 -
learning_rate: 2.5000e-05
Epoch 12/20
5/5          14s 3s/step -
accuracy: 0.5336 - loss: 0.8589 - val_accuracy: 0.5250 - val_loss: 0.7185 -
learning_rate: 2.5000e-05
Epoch 13/20
5/5          14s 3s/step -
accuracy: 0.5051 - loss: 0.8689 - val_accuracy: 0.5250 - val_loss: 0.7165 -
learning_rate: 2.5000e-05
Epoch 14/20
5/5          14s 3s/step -
accuracy: 0.5437 - loss: 0.8528 - val_accuracy: 0.5250 - val_loss: 0.7151 -
learning_rate: 1.2500e-05
Epoch 15/20
5/5          12s 2s/step -
accuracy: 0.5268 - loss: 0.8150 - val_accuracy: 0.5250 - val_loss: 0.7147 -

```



```

learning_rate: 1.2500e-05
Epoch 16/20
5/5          22s 3s/step -
accuracy: 0.4775 - loss: 0.9822 - val_accuracy: 0.5250 - val_loss: 0.7161 -
learning_rate: 1.2500e-05
Epoch 17/20
5/5          13s 3s/step -
accuracy: 0.5513 - loss: 0.8773 - val_accuracy: 0.5250 - val_loss: 0.7167 -
learning_rate: 6.2500e-06
Epoch 18/20
5/5          20s 3s/step -
accuracy: 0.5411 - loss: 0.8141 - val_accuracy: 0.5250 - val_loss: 0.7173 -
learning_rate: 6.2500e-06
Epoch 19/20
5/5          11s 2s/step -
accuracy: 0.6277 - loss: 0.7813 - val_accuracy: 0.5250 - val_loss: 0.7172 -
learning_rate: 6.2500e-06
Epoch 20/20
5/5          11s 2s/step -
accuracy: 0.4694 - loss: 0.9221 - val_accuracy: 0.5250 - val_loss: 0.7173 -
learning_rate: 3.1250e-06
2/2          1s 237ms/step -
accuracy: 0.5375 - loss: 0.7127
Test accuracy: 0.525

```

4.4 12 - Confusion Matrix Hard (disco and pop)

```

[14]: import seaborn as sns
      # from sklearn.metrics import confusion
      import numpy as NP
      from sklearn.metrics import confusion_matrix

      cnn_preds = np.argmax(model.predict(X_test), axis=1)
      cnn_cm = confusion_matrix(y_test, cnn_preds)

      # Plot the confusion matrix
      plt.figure(figsize=(8, 6))
      sns.heatmap(cnn_cm, annot=True, fmt="d", cmap="Blues", xticklabels=GENRES,
                  yticklabels=GENRES)
      plt.title("CNN Confusion Matrix")
      plt.xlabel("Predicted")
      plt.ylabel("True")
      plt.show()

```

WARNING:tensorflow:5 out of the last 10 calls to <function TensorFlowTrainer.make_predict_function.<locals>.one_step_on_data_distributed at 0x7fca884a6340> triggered tf.function retracing. Tracing is expensive and the excessive number of tracings could be due to (1) creating @tf.function

repeatedly in a loop, (2) passing tensors with different shapes, (3) passing Python objects instead of tensors. For (1), please define your `@tf.function` outside of the loop. For (2), `@tf.function` has `reduce_retracing=True` option that can avoid unnecessary retracing. For (3), please refer to https://www.tensorflow.org/guide/function#controlling_retracing and https://www.tensorflow.org/api_docs/python/tf/function for more details.


1/2

1s

1s/stepWARNING:tensorflow:6 out of the last 11 calls to <function TensorFlowTrainer.make_predict_function.<locals>.one_step_on_data_distributed at 0x7fca884a6340> triggered tf.function retracing. Tracing is expensive and the excessive number of tracings could be due to (1) creating `@tf.function` repeatedly in a loop, (2) passing tensors with different shapes, (3) passing Python objects instead of tensors. For (1), please define your `@tf.function` outside of the loop. For (2), `@tf.function` has `reduce_retracing=True` option that can avoid unnecessary retracing. For (3), please refer to https://www.tensorflow.org/guide/function#controlling_retracing and https://www.tensorflow.org/api_docs/python/tf/function for more details.

2/2

2s 990ms/step



4.5 13 - Limited Genres Medium (5 random)

```
[15]: import os
import numpy as np
import tensorflow as tf
from sklearn.model_selection import train_test_split
import random

# Augmentation function
def augment_image(image):
    image = tf.image.random_flip_left_right(image)
    image = tf.image.random_brightness(image, max_delta=0.1)
    image = tf.image.random_contrast(image, 0.8, 1.2)
    return image

# Define the genres and file paths
GENRES = ['blues', 'classical', 'country', 'disco', 'hiphop', 'jazz', 'metal', 'pop', 'reggae', 'rock']
GENRES = random.sample(GENRES, 5)
print(GENRES)

FILE_PATH = os.path.join('Data', 'mel_spectrograms', 'mel_spectrogram_32')
X = []
y = []

GENRE_TO_INDEX = {genre: index for index, genre in enumerate(GENRES)}

# Loop through the genres and load the images with augmentation
for genre in GENRES:
    genre_dir = os.path.join(FILE_PATH, genre)
    print(f"Going through {genre}")
    for file in os.listdir(genre_dir):
        image = tf.io.read_file(os.path.join(genre_dir, file))
        image = tf.image.decode_png(image, channels=1)
        image = tf.image.convert_image_dtype(image, tf.float32)
        image = tf.image.resize(image, [256, 256]) # Resize to 256x256
        image = augment_image(image) # Apply augmentation
        image = image.numpy() # Convert to numpy array
        X.append(image)
        y.append(GENRE_TO_INDEX[genre])

# Convert lists to numpy arrays
X = np.array(X)
y = np.array(y)

# Split the data into training and testing sets
```

```

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
↳random_state=42)

from tensorflow.keras import models
from tensorflow.keras.layers import Conv2D, MaxPooling2D, Flatten, Dense,
↳Dropout, Normalization

model = models.Sequential([
    Conv2D(32, (3, 3), activation='relu', input_shape=(256, 256, 1)),
    Normalization(),
    MaxPooling2D((2, 2)),

    Conv2D(64, (3, 3), activation='relu'),
    Normalization(),
    MaxPooling2D((2, 2)),

    Conv2D(128, (3, 3), activation='relu'),
    Normalization(),
    MaxPooling2D((2, 2)),

    Conv2D(256, (3, 3), activation='relu'),
    Normalization(),
    MaxPooling2D((2, 2)),

    Flatten(),

    Dense(512, activation='relu'),
    Dropout(0.5),

    Dense(256, activation='relu'),
    Dropout(0.5),

    Dense(128, activation='relu'),
    Dense(10, activation='softmax')
])

from tensorflow.keras.optimizers import Adam
from tensorflow.keras.callbacks import ReduceLROnPlateau

model.compile(optimizer=Adam(learning_rate=0.0001),
↳loss='sparse_categorical_crossentropy', metrics=['accuracy'])
reduce_lr = ReduceLROnPlateau(monitor='val_loss', factor=0.5, patience=3,
↳min_lr=1e-6)

model.fit(X_train, y_train, epochs=20, validation_data=(X_test, y_test),
↳batch_size=32, callbacks=[reduce_lr])

```

```
evaluation = model.evaluate(X_test, y_test)
print(f"Test accuracy: {evaluation[1]:.3f}")
```

```
['hiphop', 'blues', 'classical', 'jazz', 'pop']
```

```
Going through hiphop
```

```
Going through blues
```

```
Going through classical
```

```
Going through jazz
```

```
Going through pop
```

```
/opt/conda/lib/python3.12/site-
```

```
packages/keras/src/layers/convolutional/base_conv.py:107: UserWarning: Do not
pass an `input_shape`/`input_dim` argument to a layer. When using Sequential
models, prefer using an `Input(shape)` object as the first layer in the model
instead.
```

```
super().__init__(activity_regularizer=activity_regularizer, **kwargs)
```

```
Epoch 1/20
```

```
13/13          44s 3s/step -
```

```
accuracy: 0.2146 - loss: 2.1531 - val_accuracy: 0.1400 - val_loss: 1.8342 -
```

```
learning_rate: 1.0000e-04
```

```
Epoch 2/20
```

```
13/13          34s 2s/step -
```

```
accuracy: 0.1739 - loss: 1.9244 - val_accuracy: 0.1400 - val_loss: 1.7228 -
```

```
learning_rate: 1.0000e-04
```

```
Epoch 3/20
```

```
13/13          29s 2s/step -
```

```
accuracy: 0.2617 - loss: 1.8506 - val_accuracy: 0.1700 - val_loss: 1.7450 -
```

```
learning_rate: 1.0000e-04
```

```
Epoch 4/20
```

```
13/13          28s 2s/step -
```

```
accuracy: 0.2629 - loss: 1.8225 - val_accuracy: 0.1400 - val_loss: 1.6690 -
```

```
learning_rate: 1.0000e-04
```

```
Epoch 5/20
```

```
13/13          33s 3s/step -
```

```
accuracy: 0.2886 - loss: 1.7340 - val_accuracy: 0.1400 - val_loss: 1.6657 -
```

```
learning_rate: 1.0000e-04
```

```
Epoch 6/20
```

```
13/13          37s 2s/step -
```

```
accuracy: 0.2870 - loss: 1.6985 - val_accuracy: 0.1500 - val_loss: 1.6433 -
```

```
learning_rate: 1.0000e-04
```

```
Epoch 7/20
```

```
13/13          32s 2s/step -
```

```
accuracy: 0.2648 - loss: 1.6549 - val_accuracy: 0.1500 - val_loss: 1.6541 -
```

```
learning_rate: 1.0000e-04
```

```
Epoch 8/20
```

```
13/13          43s 3s/step -
```

```
accuracy: 0.2876 - loss: 1.7098 - val_accuracy: 0.2500 - val_loss: 1.6407 -
```

```
learning_rate: 1.0000e-04
```

Epoch 9/20
13/13 40s 3s/step -
accuracy: 0.2835 - loss: 1.6349 - val_accuracy: 0.4000 - val_loss: 1.5025 -
learning_rate: 1.0000e-04

Epoch 10/20
13/13 46s 3s/step -
accuracy: 0.2688 - loss: 1.5492 - val_accuracy: 0.4000 - val_loss: 1.5402 -
learning_rate: 1.0000e-04

Epoch 11/20
13/13 41s 3s/step -
accuracy: 0.3915 - loss: 1.4347 - val_accuracy: 0.4400 - val_loss: 1.3902 -
learning_rate: 1.0000e-04

Epoch 12/20
13/13 40s 3s/step -
accuracy: 0.3615 - loss: 1.4412 - val_accuracy: 0.4700 - val_loss: 1.4123 -
learning_rate: 1.0000e-04

Epoch 13/20
13/13 34s 3s/step -
accuracy: 0.3856 - loss: 1.3626 - val_accuracy: 0.4600 - val_loss: 1.2885 -
learning_rate: 1.0000e-04

Epoch 14/20
13/13 46s 3s/step -
accuracy: 0.4719 - loss: 1.2948 - val_accuracy: 0.5800 - val_loss: 1.2279 -
learning_rate: 1.0000e-04

Epoch 15/20
13/13 39s 3s/step -
accuracy: 0.4676 - loss: 1.2074 - val_accuracy: 0.5100 - val_loss: 1.1995 -
learning_rate: 1.0000e-04

Epoch 16/20
13/13 32s 2s/step -
accuracy: 0.5400 - loss: 1.1320 - val_accuracy: 0.5300 - val_loss: 1.1308 -
learning_rate: 1.0000e-04

Epoch 17/20
13/13 35s 3s/step -
accuracy: 0.5496 - loss: 1.0961 - val_accuracy: 0.5900 - val_loss: 1.0614 -
learning_rate: 1.0000e-04

Epoch 18/20
13/13 43s 3s/step -
accuracy: 0.5570 - loss: 1.0951 - val_accuracy: 0.6500 - val_loss: 1.0640 -
learning_rate: 1.0000e-04

Epoch 19/20
13/13 32s 2s/step -
accuracy: 0.6150 - loss: 0.9955 - val_accuracy: 0.5600 - val_loss: 1.0584 -
learning_rate: 1.0000e-04

Epoch 20/20
13/13 48s 3s/step -
accuracy: 0.6067 - loss: 1.0018 - val_accuracy: 0.4900 - val_loss: 1.1657 -
learning_rate: 1.0000e-04

4/4 2s 416ms/step -
accuracy: 0.4887 - loss: 1.1907
Test accuracy: 0.490

4.6 14 - Confusion Matrix Medium (5 random)

```
[16]: import seaborn as sns
      # from sklearn.metrics import confusion
      import numpy as NP
      from sklearn.metrics import confusion_matrix

      cnn_preds = np.argmax(model.predict(X_test), axis=1)
      cnn_cm = confusion_matrix(y_test, cnn_preds)

      # Plot the confusion matrix
      plt.figure(figsize=(8, 6))
      sns.heatmap(cnn_cm, annot=True, fmt="d", cmap="Blues", xticklabels=GENRES,
                  yticklabels=GENRES)
      plt.title("CNN Confusion Matrix")
      plt.xlabel("Predicted")
      plt.ylabel("True")
      plt.show()
```

4/4 2s 460ms/step

