

# Spectrogram-Only (3 secs) CNN

March 21, 2025

```
[1]: !pip install librosa
```

```
Requirement already satisfied: librosa in /opt/conda/lib/python3.12/site-  
packages (0.11.0)  
Requirement already satisfied: audioread>=2.1.9 in  
/opt/conda/lib/python3.12/site-packages (from librosa) (3.0.1)  
Requirement already satisfied: numba>=0.51.0 in /opt/conda/lib/python3.12/site-  
packages (from librosa) (0.60.0)  
Requirement already satisfied: numpy>=1.22.3 in /opt/conda/lib/python3.12/site-  
packages (from librosa) (2.0.2)  
Requirement already satisfied: scipy>=1.6.0 in /opt/conda/lib/python3.12/site-  
packages (from librosa) (1.14.1)  
Requirement already satisfied: scikit-learn>=1.1.0 in  
/opt/conda/lib/python3.12/site-packages (from librosa) (1.5.2)  
Requirement already satisfied: joblib>=1.0 in /opt/conda/lib/python3.12/site-  
packages (from librosa) (1.4.2)  
Requirement already satisfied: decorator>=4.3.0 in  
/opt/conda/lib/python3.12/site-packages (from librosa) (5.1.1)  
Requirement already satisfied: soundfile>=0.12.1 in  
/opt/conda/lib/python3.12/site-packages (from librosa) (0.13.1)  
Requirement already satisfied: pooch>=1.1 in /opt/conda/lib/python3.12/site-  
packages (from librosa) (1.8.2)  
Requirement already satisfied: soxr>=0.3.2 in /opt/conda/lib/python3.12/site-  
packages (from librosa) (0.5.0.post1)  
Requirement already satisfied: typing_extensions>=4.1.1 in  
/opt/conda/lib/python3.12/site-packages (from librosa) (4.12.2)  
Requirement already satisfied: lazy_loader>=0.1 in  
/opt/conda/lib/python3.12/site-packages (from librosa) (0.4)  
Requirement already satisfied: msgpack>=1.0 in /opt/conda/lib/python3.12/site-  
packages (from librosa) (1.1.0)  
Requirement already satisfied: packaging in /opt/conda/lib/python3.12/site-  
packages (from lazy_loader>=0.1->librosa) (24.1)  
Requirement already satisfied: llvmlite<0.44,>=0.43.0dev0 in  
/opt/conda/lib/python3.12/site-packages (from numba>=0.51.0->librosa) (0.43.0)  
Requirement already satisfied: platformdirs>=2.5.0 in  
/opt/conda/lib/python3.12/site-packages (from pooch>=1.1->librosa) (4.3.6)  
Requirement already satisfied: requests>=2.19.0 in  
/opt/conda/lib/python3.12/site-packages (from pooch>=1.1->librosa) (2.32.3)
```

Requirement already satisfied: threadpoolctl>=3.1.0 in /opt/conda/lib/python3.12/site-packages (from scikit-learn>=1.1.0->librosa) (3.5.0)  
Requirement already satisfied: cffi>=1.0 in /opt/conda/lib/python3.12/site-packages (from soundfile>=0.12.1->librosa) (1.17.1)  
Requirement already satisfied: pycparser in /opt/conda/lib/python3.12/site-packages (from cffi>=1.0->soundfile>=0.12.1->librosa) (2.22)  
Requirement already satisfied: charset-normalizer<4,>=2 in /opt/conda/lib/python3.12/site-packages (from requests>=2.19.0->pooch>=1.1->librosa) (3.4.0)  
Requirement already satisfied: idna<4,>=2.5 in /opt/conda/lib/python3.12/site-packages (from requests>=2.19.0->pooch>=1.1->librosa) (3.10)  
Requirement already satisfied: urllib3<3,>=1.21.1 in /opt/conda/lib/python3.12/site-packages (from requests>=2.19.0->pooch>=1.1->librosa) (2.2.3)  
Requirement already satisfied: certifi>=2017.4.17 in /opt/conda/lib/python3.12/site-packages (from requests>=2.19.0->pooch>=1.1->librosa) (2024.8.30)

## 1 CNN for Spectrogram (3 secs)

### 1.1 1 - All the imports

```
[2]: import os
import numpy as np
from sklearn.model_selection import train_test_split
import tensorflow as tf
```

2025-03-21 00:44:38.874410: I tensorflow/core/platform/cpu\_feature\_guard.cc:210] This TensorFlow binary is optimized to use available CPU instructions in performance-critical operations.  
To enable the following instructions: AVX2 FMA, in other operations, rebuild TensorFlow with the appropriate compiler flags.

### 1.2 2 - Put the data within the model

```
[3]: # Import a single image and save it to be read by the model

image = os.path.join('blues.00000.png')

# Load the image
image = tf.io.read_file(image)

# Convert to a numpy array
image = tf.image.decode_png(image, channels=1)
image = tf.image.convert_image_dtype(image, tf.float32)
image = tf.image.resize(image, [256, 256])
image = image.numpy()
```

## 2 3 - Create the model

```
[4]: from tensorflow.keras import models
      from tensorflow.keras.layers import Conv2D, MaxPooling2D, Flatten, Dense, \
      ↪Dropout, BatchNormalization

model = models.Sequential([
    Conv2D(32, (3, 3), activation='relu', input_shape=(256, 256, 1)),
    BatchNormalization(),
    MaxPooling2D((2, 2)),

    Conv2D(64, (3, 3), activation='relu'),
    BatchNormalization(),
    MaxPooling2D((2, 2)),

    Conv2D(128, (3, 3), activation='relu'),
    BatchNormalization(),
    MaxPooling2D((2, 2)),

    Conv2D(256, (3, 3), activation='relu'),
    BatchNormalization(),
    MaxPooling2D((2, 2)),

    Flatten(),

    Dense(512, activation='relu'),
    Dropout(0.5),

    Dense(256, activation='relu'),
    Dropout(0.5),

    Dense(128, activation='relu'),
    Dense(10, activation='softmax')
])
```

```
/opt/conda/lib/python3.12/site-
packages/keras/src/layers/convolutional/base_conv.py:107: UserWarning: Do not
pass an `input_shape`/`input_dim` argument to a layer. When using Sequential
models, prefer using an `Input(shape)` object as the first layer in the model
instead.
```

```
    super().__init__(activity_regularizer=activity_regularizer, **kwargs)
```

```
[5]: model.summary()
```

```
Model: "sequential"
```

Layer (type)

Output Shape

Param #

conv2d (Conv2D)	(None, 254, 254, 32)	320
batch_normalization (BatchNormalization)	(None, 254, 254, 32)	128
max_pooling2d (MaxPooling2D)	(None, 127, 127, 32)	0
conv2d_1 (Conv2D)	(None, 125, 125, 64)	18,496
batch_normalization_1 (BatchNormalization)	(None, 125, 125, 64)	256
max_pooling2d_1 (MaxPooling2D)	(None, 62, 62, 64)	0
conv2d_2 (Conv2D)	(None, 60, 60, 128)	73,856
batch_normalization_2 (BatchNormalization)	(None, 60, 60, 128)	512
max_pooling2d_2 (MaxPooling2D)	(None, 30, 30, 128)	0
conv2d_3 (Conv2D)	(None, 28, 28, 256)	295,168
batch_normalization_3 (BatchNormalization)	(None, 28, 28, 256)	1,024
max_pooling2d_3 (MaxPooling2D)	(None, 14, 14, 256)	0
flatten (Flatten)	(None, 50176)	0
dense (Dense)	(None, 512)	25,690,624
dropout (Dropout)	(None, 512)	0
dense_1 (Dense)	(None, 256)	131,328
dropout_1 (Dropout)	(None, 256)	0
dense_2 (Dense)	(None, 128)	32,896
dense_3 (Dense)	(None, 10)	1,290

Total params: 26,245,898 (100.12 MB)

Trainable params: 26,244,938 (100.12 MB)

Non-trainable params: 960 (3.75 KB)

### 3 4 - Load the images

```
[6]: import tensorflow as tf
import os
import numpy as np
from sklearn.model_selection import train_test_split

GENRES = ['blues', 'classical', 'country', 'disco', 'hiphop', 'jazz', 'metal', '
↳pop', 'reggae', 'rock']
FILE_PATH = os.path.join('Data', 'spectrograms (3 secs)', 'spectrogram_256')
X = []
y = []

GENRE_TO_INDEX = {genre: index for index, genre in enumerate(GENRES)}

# Define the augmentation function
def augment_image(image):
    image = tf.image.random_flip_left_right(image)
    image = tf.image.random_brightness(image, max_delta=0.1)
    image = tf.image.random_contrast(image, 0.8, 1.2)
    return image

for genre in GENRES:
    genre_dir = os.path.join(FILE_PATH, genre)
    print(f"Going through {genre}")
    for file in os.listdir(genre_dir):
        try:
            image = tf.io.read_file(os.path.join(genre_dir, file))
            image = tf.image.decode_png(image, channels=1)
            image = tf.image.convert_image_dtype(image, tf.float32)
            image = tf.image.resize(image, [256, 256])

            # Apply the augmentation
            image = augment_image(image)

            image = image.numpy() # Convert to numpy array for further
↳processing
            X.append(image)
            y.append(GENRE_TO_INDEX[genre])
        except:
            continue
```

```

X = np.array(X)
y = np.array(y)

# Split the data
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
↳random_state=42)

```

Going through blues  
 Going through classical  
 Going through country  
 Going through disco  
 Going through hiphop  
 Going through jazz  
 Going through metal  
 Going through pop  
 Going through reggae  
 Going through rock

2025-03-21 00:57:56.983206: W tensorflow/core/framework/op\_kernel.cc:1841]  
 OP\_REQUIRES failed at whole\_file\_read\_ops.cc:116 : FAILED\_PRECONDITION:  
 Data/spectrograms (3 secs)/spectrogram\_256/rock/.ipynb\_checkpoints; Is a  
 directory  
 2025-03-21 00:57:56.988155: I tensorflow/core/framework/local\_rendezvous.cc:405]  
 Local rendezvous is aborting with status: FAILED\_PRECONDITION: Data/spectrograms  
 (3 secs)/spectrogram\_256/rock/.ipynb\_checkpoints; Is a directory


```
[7]: X.shape, y.shape
```

```
[7]: ((10000, 256, 256, 1), (10000,))
```

```

[8]: # Show image as a sanity check
import matplotlib.pyplot as plt
plt.imshow(X_train[7999].reshape(256, 256))
plt.show()

```



```
[9]: from tensorflow.keras.optimizers import Adam
      from tensorflow.keras.callbacks import ReduceLROnPlateau

      model.compile(optimizer=Adam(learning_rate=0.0001),
                    ↪loss='sparse_categorical_crossentropy', metrics=['accuracy'])

      reduce_lr = ReduceLROnPlateau(monitor='val_loss', factor=0.5, patience=3,
                                     ↪min_lr=1e-6)
```

```
[10]: model.fit(X_train, y_train, epochs=20, validation_data=(X_test, y_test),
                ↪batch_size=32, callbacks=[reduce_lr])
```

Epoch 1/20

250/250 1381s 5s/step -

accuracy: 0.1601 - loss: 3.1593 - val\_accuracy: 0.1010 - val\_loss: 5.0143 -

learning\_rate: 1.0000e-04

Epoch 2/20

250/250 1189s 5s/step -

accuracy: 0.2633 - loss: 2.0170 - val\_accuracy: 0.2210 - val\_loss: 2.5034 -

learning\_rate: 1.0000e-04

Epoch 3/20

250/250 964s 4s/step -

accuracy: 0.2974 - loss: 1.9189 - val\_accuracy: 0.3580 - val\_loss: 1.7114 -  
 learning\_rate: 1.0000e-04  
 Epoch 4/20  
 250/250 784s 3s/step -  
 accuracy: 0.3523 - loss: 1.7837 - val\_accuracy: 0.4620 - val\_loss: 1.4956 -  
 learning\_rate: 1.0000e-04  
 Epoch 5/20  
 250/250 801s 3s/step -  
 accuracy: 0.3829 - loss: 1.6866 - val\_accuracy: 0.4935 - val\_loss: 1.4203 -  
 learning\_rate: 1.0000e-04  
 Epoch 6/20  
 250/250 791s 3s/step -  
 accuracy: 0.4342 - loss: 1.5779 - val\_accuracy: 0.5180 - val\_loss: 1.3586 -  
 learning\_rate: 1.0000e-04  
 Epoch 7/20  
 250/250 795s 3s/step -  
 accuracy: 0.4765 - loss: 1.4738 - val\_accuracy: 0.6180 - val\_loss: 1.2474 -  
 learning\_rate: 1.0000e-04  
 Epoch 8/20  
 250/250 637s 3s/step -  
 accuracy: 0.5191 - loss: 1.3491 - val\_accuracy: 0.6615 - val\_loss: 1.1249 -  
 learning\_rate: 1.0000e-04  
 Epoch 9/20  
 250/250 631s 3s/step -  
 accuracy: 0.5544 - loss: 1.2650 - val\_accuracy: 0.6140 - val\_loss: 1.1917 -  
 learning\_rate: 1.0000e-04  
 Epoch 10/20  
 250/250 641s 3s/step -  
 accuracy: 0.5804 - loss: 1.1998 - val\_accuracy: 0.6580 - val\_loss: 1.0853 -  
 learning\_rate: 1.0000e-04  
 Epoch 11/20  
 250/250 639s 3s/step -  
 accuracy: 0.6206 - loss: 1.0845 - val\_accuracy: 0.6965 - val\_loss: 1.0171 -  
 learning\_rate: 1.0000e-04  
 Epoch 12/20  
 250/250 621s 2s/step -  
 accuracy: 0.6481 - loss: 1.0130 - val\_accuracy: 0.6905 - val\_loss: 1.0017 -  
 learning\_rate: 1.0000e-04  
 Epoch 13/20  
 250/250 629s 3s/step -  
 accuracy: 0.6793 - loss: 0.9440 - val\_accuracy: 0.6965 - val\_loss: 1.0273 -  
 learning\_rate: 1.0000e-04  
 Epoch 14/20  
 250/250 621s 2s/step -  
 accuracy: 0.7040 - loss: 0.8534 - val\_accuracy: 0.6575 - val\_loss: 1.1556 -  
 learning\_rate: 1.0000e-04  
 Epoch 15/20  
 250/250 634s 3s/step -



```

accuracy: 0.7099 - loss: 0.8308 - val_accuracy: 0.5790 - val_loss: 1.4206 -
learning_rate: 1.0000e-04
Epoch 16/20
250/250          618s 2s/step -
accuracy: 0.7318 - loss: 0.7707 - val_accuracy: 0.7540 - val_loss: 0.8052 -
learning_rate: 5.0000e-05
Epoch 17/20
250/250          615s 2s/step -
accuracy: 0.7702 - loss: 0.6519 - val_accuracy: 0.7765 - val_loss: 0.7779 -
learning_rate: 5.0000e-05
Epoch 18/20
250/250          618s 2s/step -
accuracy: 0.7737 - loss: 0.6390 - val_accuracy: 0.7645 - val_loss: 0.7963 -
learning_rate: 5.0000e-05
Epoch 19/20
250/250          627s 3s/step -
accuracy: 0.7821 - loss: 0.6070 - val_accuracy: 0.7655 - val_loss: 0.7982 -
learning_rate: 5.0000e-05
Epoch 20/20
250/250          620s 2s/step -
accuracy: 0.7976 - loss: 0.5731 - val_accuracy: 0.7640 - val_loss: 0.8039 -
learning_rate: 5.0000e-05

```

[10]: <keras.src.callbacks.history.History at 0x7f4d60618b90>

```

[11]: evaluation = model.evaluate(X_test, y_test)
print(f"Test accuracy: {evaluation[1]:.3f}")

```

```

63/63          19s 304ms/step -
accuracy: 0.7480 - loss: 0.8153
Test accuracy: 0.764

```

## 4 Apply the confusion matrix after the model

```

[12]: import seaborn as sns
# from sklearn.metrics import confusion
import numpy as NP
from sklearn.metrics import confusion_matrix

cnn_preds = np.argmax(model.predict(X_test), axis=1)
cnn_cm = confusion_matrix(y_test, cnn_preds)


# Plot the confusion matrix
plt.figure(figsize=(8, 6))
sns.heatmap(cnn_cm, annot=True, fmt="d", cmap="Blues", xticklabels=GENRES,
            yticklabels=GENRES)
plt.title("CNN Confusion Matrix")

```

```
plt.xlabel("Predicted")
plt.ylabel("True")
plt.show()
```

63/63

20s 309ms/step



```
[13]: X.shape, y.shape
```

```
[13]: ((10000, 256, 256, 1), (10000,))
```

#### 4.1 9 - Limited Genres Easy (metal and classical)

```
[14]: import tensorflow as tf
import os
import numpy as np
from sklearn.model_selection import train_test_split

GENRES = ['classical', 'metal']
```

```

FILE_PATH = os.path.join('Data', 'spectrograms (3 secs)', 'spectrogram_256')
X = []
y = []

GENRE_TO_INDEX = {genre: index for index, genre in enumerate(GENRES)}

# Define the augmentation function
def augment_image(image):
    image = tf.image.random_flip_left_right(image)
    image = tf.image.random_brightness(image, max_delta=0.1)
    image = tf.image.random_contrast(image, 0.8, 1.2)
    return image

for genre in GENRES:
    genre_dir = os.path.join(FILE_PATH, genre)
    print(f"Going through {genre}")
    for file in os.listdir(genre_dir):
        image = tf.io.read_file(os.path.join(genre_dir, file))
        image = tf.image.decode_png(image, channels=1)
        image = tf.image.convert_image_dtype(image, tf.float32)
        image = tf.image.resize(image, [256, 256])

        # Apply the augmentation
        image = augment_image(image)

        image = image.numpy() # Convert to numpy array for further processing
        X.append(image)
        y.append(GENRE_TO_INDEX[genre])

X = np.array(X)
y = np.array(y)

# Split the data
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
    ↪random_state=42)

from tensorflow.keras import models
from tensorflow.keras.layers import Conv2D, MaxPooling2D, Flatten, Dense,
    ↪Dropout, Normalization

model = models.Sequential([
    Conv2D(32, (3, 3), activation='relu', input_shape=(256, 256, 1)),
    Normalization(),
    MaxPooling2D((2, 2)),

    Conv2D(64, (3, 3), activation='relu'),
    Normalization(),
    MaxPooling2D((2, 2)),

```

```

Conv2D(128, (3, 3), activation='relu'),
Normalization(),
MaxPooling2D((2, 2)),

Conv2D(256, (3, 3), activation='relu'),
Normalization(),
MaxPooling2D((2, 2)),

Flatten(),

Dense(512, activation='relu'),
Dropout(0.5),

Dense(256, activation='relu'),
Dropout(0.5),

Dense(128, activation='relu'),
Dense(10, activation='softmax')
])

from tensorflow.keras.optimizers import Adam
from tensorflow.keras.callbacks import ReduceLROnPlateau

model.compile(optimizer=Adam(learning_rate=0.0001),
    ↳loss='sparse_categorical_crossentropy', metrics=['accuracy'])
reduce_lr = ReduceLROnPlateau(monitor='val_loss', factor=0.5, patience=3,
    ↳min_lr=1e-6)

model.fit(X_train, y_train, epochs=20, validation_data=(X_test, y_test),
    ↳batch_size=32, callbacks=[reduce_lr])

evaluation = model.evaluate(X_test, y_test)
print(f"Test accuracy: {evaluation[1]:.3f}")

```

Going through classical

Going through metal

/opt/conda/lib/python3.12/site-

packages/keras/src/layers/convolutional/base\_conv.py:107: UserWarning: Do not pass an `input\_shape`/`input\_dim` argument to a layer. When using Sequential models, prefer using an `Input(shape)` object as the first layer in the model instead.

```
super().__init__(activity_regularizer=activity_regularizer, **kwargs)
```

Epoch 1/20

50/50 76s 1s/step -

accuracy: 0.4420 - loss: 1.6141 - val\_accuracy: 0.8700 - val\_loss: 0.3231 -

```

learning_rate: 1.0000e-04
Epoch 2/20
50/50          70s 1s/step -
accuracy: 0.8854 - loss: 0.3115 - val_accuracy: 0.9500 - val_loss: 0.1281 -
learning_rate: 1.0000e-04
Epoch 3/20
50/50          70s 1s/step -
accuracy: 0.9582 - loss: 0.1412 - val_accuracy: 0.9625 - val_loss: 0.1082 -
learning_rate: 1.0000e-04
Epoch 4/20
50/50          69s 1s/step -
accuracy: 0.9631 - loss: 0.1169 - val_accuracy: 0.9475 - val_loss: 0.1402 -
learning_rate: 1.0000e-04
Epoch 5/20
50/50          69s 1s/step -
accuracy: 0.9697 - loss: 0.0994 - val_accuracy: 0.9875 - val_loss: 0.0373 -
learning_rate: 1.0000e-04
Epoch 6/20
50/50          71s 1s/step -
accuracy: 0.9785 - loss: 0.0791 - val_accuracy: 0.9900 - val_loss: 0.0332 -
learning_rate: 1.0000e-04
Epoch 7/20
50/50          71s 1s/step -
accuracy: 0.9821 - loss: 0.0587 - val_accuracy: 0.9750 - val_loss: 0.0527 -
learning_rate: 1.0000e-04
Epoch 8/20
50/50          71s 1s/step -
accuracy: 0.9730 - loss: 0.0675 - val_accuracy: 0.9950 - val_loss: 0.0226 -
learning_rate: 1.0000e-04
Epoch 9/20
50/50          75s 1s/step -
accuracy: 0.9894 - loss: 0.0408 - val_accuracy: 0.9950 - val_loss: 0.0201 -
learning_rate: 1.0000e-04
Epoch 10/20
50/50          71s 1s/step -
accuracy: 0.9921 - loss: 0.0352 - val_accuracy: 0.9925 - val_loss: 0.0176 -
learning_rate: 1.0000e-04
Epoch 11/20
50/50          72s 1s/step -
accuracy: 0.9888 - loss: 0.0411 - val_accuracy: 0.9950 - val_loss: 0.0168 -
learning_rate: 1.0000e-04
Epoch 12/20
50/50          71s 1s/step -
accuracy: 0.9886 - loss: 0.0326 - val_accuracy: 0.9925 - val_loss: 0.0193 -
learning_rate: 1.0000e-04
Epoch 13/20
50/50          72s 1s/step -
accuracy: 0.9943 - loss: 0.0199 - val_accuracy: 0.9975 - val_loss: 0.0071 -

```

```

learning_rate: 1.0000e-04
Epoch 14/20
50/50          72s 1s/step -
accuracy: 0.9981 - loss: 0.0122 - val_accuracy: 0.9975 - val_loss: 0.0073 -
learning_rate: 1.0000e-04
Epoch 15/20
50/50          57s 1s/step -
accuracy: 0.9953 - loss: 0.0206 - val_accuracy: 1.0000 - val_loss: 0.0042 -
learning_rate: 1.0000e-04
Epoch 16/20
50/50          57s 1s/step -
accuracy: 0.9930 - loss: 0.0226 - val_accuracy: 0.9925 - val_loss: 0.0180 -
learning_rate: 1.0000e-04
Epoch 17/20
50/50          57s 1s/step -
accuracy: 0.9937 - loss: 0.0228 - val_accuracy: 1.0000 - val_loss: 0.0033 -
learning_rate: 1.0000e-04
Epoch 18/20
50/50          57s 1s/step -
accuracy: 0.9971 - loss: 0.0107 - val_accuracy: 0.9975 - val_loss: 0.0068 -
learning_rate: 1.0000e-04
Epoch 19/20
50/50          57s 1s/step -
accuracy: 0.9972 - loss: 0.0191 - val_accuracy: 0.9950 - val_loss: 0.0074 -
learning_rate: 1.0000e-04
Epoch 20/20
50/50          57s 1s/step -
accuracy: 0.9983 - loss: 0.0128 - val_accuracy: 0.9950 - val_loss: 0.0119 -
learning_rate: 1.0000e-04
13/13          3s 223ms/step -
accuracy: 0.9955 - loss: 0.0133
Test accuracy: 0.995

```

## 4.2 10 - Confusion Matrix Easy (classical and metal)

```

[15]: import seaborn as sns
      # from sklearn.metrics import confusion
      import numpy as NP
      from sklearn.metrics import confusion_matrix

      cnn_preds = np.argmax(model.predict(X_test), axis=1)
      cnn_cm = confusion_matrix(y_test, cnn_preds)


      # Plot the confusion matrix
      plt.figure(figsize=(8, 6))
      sns.heatmap(cnn_cm, annot=True, fmt="d", cmap="Blues", xticklabels=GENRES,
                  yticklabels=GENRES)

```

```
plt.title("CNN Confusion Matrix")
plt.xlabel("Predicted")
plt.ylabel("True")
plt.show()
```

13/13

3s 240ms/step



#### 4.3 11 - Limited genres Hard (disco and pop)

```
[16]: import tensorflow as tf
import os
import numpy as np
from sklearn.model_selection import train_test_split

GENRES = ['disco', 'pop']
FILE_PATH = os.path.join('Data', 'spectrograms (3 secs)', 'spectrogram_256')
X = []
y = []
```

```

GENRE_TO_INDEX = {genre: index for index, genre in enumerate(GENRES)}

# Define the augmentation function
def augment_image(image):
    image = tf.image.random_flip_left_right(image)
    image = tf.image.random_brightness(image, max_delta=0.1)
    image = tf.image.random_contrast(image, 0.8, 1.2)
    return image

for genre in GENRES:
    genre_dir = os.path.join(FILE_PATH, genre)
    print(f"Going through {genre}")
    for file in os.listdir(genre_dir):
        image = tf.io.read_file(os.path.join(genre_dir, file))
        image = tf.image.decode_png(image, channels=1)
        image = tf.image.convert_image_dtype(image, tf.float32)
        image = tf.image.resize(image, [256, 256])

        # Apply the augmentation
        image = augment_image(image)

        image = image.numpy() # Convert to numpy array for further processing
        X.append(image)
        y.append(GENRE_TO_INDEX[genre])

X = np.array(X)
y = np.array(y)

# Split the data
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
    random_state=42)
from tensorflow.keras import models
from tensorflow.keras.layers import Conv2D, MaxPooling2D, Flatten, Dense,
    Dropout, Normalization

model = models.Sequential([
    Conv2D(32, (3, 3), activation='relu', input_shape=(256, 256, 1)),
    Normalization(),
    MaxPooling2D((2, 2)),

    Conv2D(64, (3, 3), activation='relu'),
    Normalization(),
    MaxPooling2D((2, 2)),

    Conv2D(128, (3, 3), activation='relu'),
    Normalization(),

```



```

MaxPooling2D((2, 2)),

Conv2D(256, (3, 3), activation='relu'),
Normalization(),
MaxPooling2D((2, 2)),

Flatten(),

Dense(512, activation='relu'),
Dropout(0.5),

Dense(256, activation='relu'),
Dropout(0.5),

Dense(128, activation='relu'),
Dense(10, activation='softmax')
])

from tensorflow.keras.optimizers import Adam
from tensorflow.keras.callbacks import ReduceLROnPlateau

model.compile(optimizer=Adam(learning_rate=0.0001),
    ↳loss='sparse_categorical_crossentropy', metrics=['accuracy'])
reduce_lr = ReduceLROnPlateau(monitor='val_loss', factor=0.5, patience=3,
    ↳min_lr=1e-6)

model.fit(X_train, y_train, epochs=20, validation_data=(X_test, y_test),
    ↳batch_size=32, callbacks=[reduce_lr])

evaluation = model.evaluate(X_test, y_test)
print(f"Test accuracy: {evaluation[1]:.3f}")

```

Going through disco

Going through pop

/opt/conda/lib/python3.12/site-packages/keras/src/layers/convolutional/base\_conv.py:107: UserWarning: Do not pass an `input\_shape`/`input\_dim` argument to a layer. When using Sequential models, prefer using an `Input(shape)` object as the first layer in the model instead.

```
super().__init__(activity_regularizer=activity_regularizer, **kwargs)
```

Epoch 1/20

50/50 59s 1s/step -

accuracy: 0.4197 - loss: 1.4163 - val\_accuracy: 0.5025 - val\_loss: 0.7680 -

learning\_rate: 1.0000e-04

Epoch 2/20

50/50 55s 1s/step -

accuracy: 0.5129 - loss: 0.8277 - val\_accuracy: 0.7900 - val\_loss: 0.4794 -  
 learning\_rate: 1.0000e-04  
 Epoch 3/20  
 50/50 55s 1s/step -  
 accuracy: 0.7707 - loss: 0.5066 - val\_accuracy: 0.8050 - val\_loss: 0.4059 -  
 learning\_rate: 1.0000e-04  
 Epoch 4/20  
 50/50 55s 1s/step -  
 accuracy: 0.7905 - loss: 0.4676 - val\_accuracy: 0.8050 - val\_loss: 0.3956 -  
 learning\_rate: 1.0000e-04  
 Epoch 5/20  
 50/50 55s 1s/step -  
 accuracy: 0.8020 - loss: 0.4264 - val\_accuracy: 0.8100 - val\_loss: 0.3481 -  
 learning\_rate: 1.0000e-04  
 Epoch 6/20  
 50/50 55s 1s/step -  
 accuracy: 0.8079 - loss: 0.4242 - val\_accuracy: 0.8125 - val\_loss: 0.3571 -  
 learning\_rate: 1.0000e-04  
 Epoch 7/20  
 50/50 54s 1s/step -  
 accuracy: 0.8150 - loss: 0.4271 - val\_accuracy: 0.8075 - val\_loss: 0.3491 -  
 learning\_rate: 1.0000e-04  
 Epoch 8/20  
 50/50 54s 1s/step -  
 accuracy: 0.8082 - loss: 0.3994 - val\_accuracy: 0.8275 - val\_loss: 0.3267 -  
 learning\_rate: 1.0000e-04  
 Epoch 9/20  
 50/50 55s 1s/step -  
 accuracy: 0.8205 - loss: 0.4061 - val\_accuracy: 0.8300 - val\_loss: 0.3107 -  
 learning\_rate: 1.0000e-04  
 Epoch 10/20  
 50/50 55s 1s/step -  
 accuracy: 0.8326 - loss: 0.3448 - val\_accuracy: 0.8325 - val\_loss: 0.3277 -  
 learning\_rate: 1.0000e-04  
 Epoch 11/20  
 50/50 55s 1s/step -  
 accuracy: 0.8143 - loss: 0.3992 - val\_accuracy: 0.8200 - val\_loss: 0.3459 -  
 learning\_rate: 1.0000e-04  
 Epoch 12/20  
 50/50 54s 1s/step -  
 accuracy: 0.8269 - loss: 0.3619 - val\_accuracy: 0.8375 - val\_loss: 0.2814 -  
 learning\_rate: 1.0000e-04  
 Epoch 13/20  
 50/50 55s 1s/step -  
 accuracy: 0.8195 - loss: 0.3675 - val\_accuracy: 0.8500 - val\_loss: 0.2907 -  
 learning\_rate: 1.0000e-04  
 Epoch 14/20  
 50/50 55s 1s/step -

```

accuracy: 0.8445 - loss: 0.3248 - val_accuracy: 0.8550 - val_loss: 0.2823 -
learning_rate: 1.0000e-04
Epoch 15/20
50/50          54s 1s/step -
accuracy: 0.8583 - loss: 0.3247 - val_accuracy: 0.8600 - val_loss: 0.2657 -
learning_rate: 1.0000e-04
Epoch 16/20
50/50          54s 1s/step -
accuracy: 0.8476 - loss: 0.3249 - val_accuracy: 0.8775 - val_loss: 0.2565 -
learning_rate: 1.0000e-04
Epoch 17/20
50/50          56s 1s/step -
accuracy: 0.8702 - loss: 0.3034 - val_accuracy: 0.8825 - val_loss: 0.2449 -
learning_rate: 1.0000e-04
Epoch 18/20
50/50          55s 1s/step -
accuracy: 0.8727 - loss: 0.2783 - val_accuracy: 0.8825 - val_loss: 0.2508 -
learning_rate: 1.0000e-04
Epoch 19/20
50/50          55s 1s/step -
accuracy: 0.8603 - loss: 0.2923 - val_accuracy: 0.9100 - val_loss: 0.2228 -
learning_rate: 1.0000e-04
Epoch 20/20
50/50          56s 1s/step -
accuracy: 0.8915 - loss: 0.2603 - val_accuracy: 0.9125 - val_loss: 0.2111 -
learning_rate: 1.0000e-04
13/13          3s 221ms/step -
accuracy: 0.9032 - loss: 0.2188
Test accuracy: 0.913

```

#### 4.4 12 - Confusion Matrix Hard (disco and pop)

```

[17]: import seaborn as sns
      # from sklearn.metrics import confusion
      import numpy as NP
      from sklearn.metrics import confusion_matrix

      cnn_preds = np.argmax(model.predict(X_test), axis=1)
      cnn_cm = confusion_matrix(y_test, cnn_preds)


      # Plot the confusion matrix
      plt.figure(figsize=(8, 6))
      sns.heatmap(cnn_cm, annot=True, fmt="d", cmap="Blues", xticklabels=GENRES,
                  yticklabels=GENRES)
      plt.title("CNN Confusion Matrix")
      plt.xlabel("Predicted")
      plt.ylabel("True")

```

```
plt.show()
```

13/13

3s 233ms/step



#### 4.5 13 - Limited Genres Medium (5 random)

```
[18]: import tensorflow as tf
import os
import numpy as np
from sklearn.model_selection import train_test_split
import random

GENRES = ['disco', 'pop']
GENRES = random.sample(GENRES, 5)
print(GENRES)
FILE_PATH = os.path.join('Data', 'spectrograms (3 secs)', 'spectrogram_256')
X = []
y = []
```

```

GENRE_TO_INDEX = {genre: index for index, genre in enumerate(GENRES)}

# Define the augmentation function
def augment_image(image):
    image = tf.image.random_flip_left_right(image)
    image = tf.image.random_brightness(image, max_delta=0.1)
    image = tf.image.random_contrast(image, 0.8, 1.2)
    return image

for genre in GENRES:
    genre_dir = os.path.join(FILE_PATH, genre)
    print(f"Going through {genre}")
    for file in os.listdir(genre_dir):
        image = tf.io.read_file(os.path.join(genre_dir, file))
        image = tf.image.decode_png(image, channels=1)
        image = tf.image.convert_image_dtype(image, tf.float32)
        image = tf.image.resize(image, [256, 256])

        # Apply the augmentation
        image = augment_image(image)

        image = image.numpy() # Convert to numpy array for further processing
        X.append(image)
        y.append(GENRE_TO_INDEX[genre])

X = np.array(X)
y = np.array(y)

# Split the data
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
    ↪random_state=42)

from tensorflow.keras import models
from tensorflow.keras.layers import Conv2D, MaxPooling2D, Flatten, Dense,
    ↪Dropout, Normalization

model = models.Sequential([
    Conv2D(32, (3, 3), activation='relu', input_shape=(256, 256, 1)),
    Normalization(),
    MaxPooling2D((2, 2)),

    Conv2D(64, (3, 3), activation='relu'),
    Normalization(),
    MaxPooling2D((2, 2)),

    Conv2D(128, (3, 3), activation='relu'),

```

```

Normalization(),
MaxPooling2D((2, 2)),

Conv2D(256, (3, 3), activation='relu'),
Normalization(),
MaxPooling2D((2, 2)),

Flatten(),

Dense(512, activation='relu'),
Dropout(0.5),

Dense(256, activation='relu'),
Dropout(0.5),

Dense(128, activation='relu'),
Dense(10, activation='softmax')
])

from tensorflow.keras.optimizers import Adam
from tensorflow.keras.callbacks import ReduceLROnPlateau

model.compile(optimizer=Adam(learning_rate=0.0001),
              ↪loss='sparse_categorical_crossentropy', metrics=['accuracy'])
reduce_lr = ReduceLROnPlateau(monitor='val_loss', factor=0.5, patience=3,
                              ↪min_lr=1e-6)

model.fit(X_train, y_train, epochs=20, validation_data=(X_test, y_test),
          ↪batch_size=32, callbacks=[reduce_lr])

evaluation = model.evaluate(X_test, y_test)
print(f"Test accuracy: {evaluation[1]:.3f}")

```

```

-----
ValueError                                Traceback (most recent call last)
Cell In[18], line 8
      5 import random
      7 GENRES = ['disco', 'pop']
----> 8 GENRES = random.sample(GENRES, 5)
      9 print(GENRES)
     10 FILE_PATH = os.path.join('Data', 'spectrograms (3 secs)',
     ↪ 'spectrogram_256')

```

```

File /opt/conda/lib/python3.12/random.py:430, in Random.sample(self, population,
↪k, counts)
     428 randbelow = self._randbelow
     429 if not 0 <= k <= n:

```

```
--> 430     raise ValueError("Sample larger than population or is negative")
      431 result = [None] * k
      432 setsize = 21          # size of a small set minus size of an empty list
```

```
ValueError: Sample larger than population or is negative
```

#### 4.6 14 - Confusion Matrix Medium (5 random)

```
[ ]: import seaborn as sns
      # from sklearn.metrics import confusion
      import numpy as NP
      from sklearn.metrics import confusion_matrix

      cnn_preds = np.argmax(model.predict(X_test), axis=1)
      cnn_cm = confusion_matrix(y_test, cnn_preds)

      # Plot the confusion matrix
      plt.figure(figsize=(8, 6))
      sns.heatmap(cnn_cm, annot=True, fmt="d", cmap="Blues", xticklabels=GENRES,
                  yticklabels=GENRES)
      plt.title("CNN Confusion Matrix")
      plt.xlabel("Predicted")
      plt.ylabel("True")
      plt.show()
```