

Tempogram-Only (3 secs) CNN

March 20, 2025

1 CNN for Tempogram (3 secs)

1.1 1 - All the imports

```
[1]: import os
import numpy as np
from sklearn.model_selection import train_test_split
import tensorflow as tf
```

2025-03-20 08:41:59.031852: I tensorflow/core/platform/cpu_feature_guard.cc:210] This TensorFlow binary is optimized to use available CPU instructions in performance-critical operations.
To enable the following instructions: AVX2 FMA, in other operations, rebuild TensorFlow with the appropriate compiler flags.

1.2 2 - Put the data within the model

```
[2]: # Import a single image and save it to be read by the model

image = os.path.join('blues.00000.png')

# Load the image
image = tf.io.read_file(image)

# Convert to a numpy array
image = tf.image.decode_png(image, channels=1)
image = tf.image.convert_image_dtype(image, tf.float32)
image = tf.image.resize(image, [256, 256])
image = image.numpy()
```

2 3 - Create the model

```
[3]: from tensorflow.keras import models
from tensorflow.keras.layers import Conv2D, MaxPooling2D, Flatten, Dense, \
↳ Dropout, BatchNormalization

model = models.Sequential([
```

```

Conv2D(32, (3, 3), activation='relu', input_shape=(256, 256, 1)),
BatchNormalization(),
MaxPooling2D((2, 2)),

Conv2D(64, (3, 3), activation='relu'),
BatchNormalization(),
MaxPooling2D((2, 2)),

Conv2D(128, (3, 3), activation='relu'),
BatchNormalization(),
MaxPooling2D((2, 2)),

Conv2D(256, (3, 3), activation='relu'),
BatchNormalization(),
MaxPooling2D((2, 2)),

Flatten(),

Dense(512, activation='relu'),
Dropout(0.5),

Dense(256, activation='relu'),
Dropout(0.5),

Dense(128, activation='relu'),
Dense(10, activation='softmax')
])

```

/opt/conda/lib/python3.12/site-packages/keras/src/layers/convolutional/base_conv.py:107: UserWarning: Do not pass an `input_shape`/`input_dim` argument to a layer. When using Sequential models, prefer using an `Input(shape)` object as the first layer in the model instead.

```
super().__init__(activity_regularizer=activity_regularizer, **kwargs)
```

```
[4]: model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 254, 254, 32)	320
batch_normalization (BatchNormalization)	(None, 254, 254, 32)	128
max_pooling2d (MaxPooling2D)	(None, 127, 127, 32)	0

conv2d_1 (Conv2D)	(None, 125, 125, 64)	18,496
batch_normalization_1 (BatchNormalization)	(None, 125, 125, 64)	256
max_pooling2d_1 (MaxPooling2D)	(None, 62, 62, 64)	0
conv2d_2 (Conv2D)	(None, 60, 60, 128)	73,856
batch_normalization_2 (BatchNormalization)	(None, 60, 60, 128)	512
max_pooling2d_2 (MaxPooling2D)	(None, 30, 30, 128)	0
conv2d_3 (Conv2D)	(None, 28, 28, 256)	295,168
batch_normalization_3 (BatchNormalization)	(None, 28, 28, 256)	1,024
max_pooling2d_3 (MaxPooling2D)	(None, 14, 14, 256)	0
flatten (Flatten)	(None, 50176)	0
dense (Dense)	(None, 512)	25,690,624
dropout (Dropout)	(None, 512)	0
dense_1 (Dense)	(None, 256)	131,328
dropout_1 (Dropout)	(None, 256)	0
dense_2 (Dense)	(None, 128)	32,896
dense_3 (Dense)	(None, 10)	1,290

Total params: 26,245,898 (100.12 MB)

Trainable params: 26,244,938 (100.12 MB)

Non-trainable params: 960 (3.75 KB)

3 4 - Load the images

```
[5]: import tensorflow as tf
import os
import numpy as np
from sklearn.model_selection import train_test_split

GENRES = ['blues', 'classical', 'country', 'disco', 'hiphop', 'jazz', 'metal', 'pop', 'reggae', 'rock']
FILE_PATH = os.path.join('Data', 'tempograms (3 secs)')
X = []
y = []

GENRE_TO_INDEX = {genre: index for index, genre in enumerate(GENRES)}

# Define the augmentation function
def augment_image(image):
    image = tf.image.random_flip_left_right(image)
    image = tf.image.random_brightness(image, max_delta=0.1)
    image = tf.image.random_contrast(image, 0.8, 1.2)
    return image

for genre in GENRES:
    genre_dir = os.path.join(FILE_PATH, genre)
    print(f"Going through {genre}")
    for file in os.listdir(genre_dir):
        image = tf.io.read_file(os.path.join(genre_dir, file))
        image = tf.image.decode_png(image, channels=1)
        image = tf.image.convert_image_dtype(image, tf.float32)
        image = tf.image.resize(image, [256, 256])

        # Apply the augmentation
        image = augment_image(image)

        image = image.numpy() # Convert to numpy array for further processing
        X.append(image)
        y.append(GENRE_TO_INDEX[genre])


X = np.array(X)
y = np.array(y)

# Split the data
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

Going through blues
Going through classical
Going through country

Going through disco
Going through hiphop
Going through jazz
Going through metal
Going through pop
Going through reggae
Going through rock

```
[6]: # Show image as a sanity check
import matplotlib.pyplot as plt
plt.imshow(X_train[22].reshape(256, 256))
plt.show()
```



```
[7]: from tensorflow.keras.optimizers import Adam
from tensorflow.keras.callbacks import ReduceLROnPlateau

model.compile(optimizer=Adam(learning_rate=0.0001),
    ↳ loss='sparse_categorical_crossentropy', metrics=['accuracy'])

reduce_lr = ReduceLROnPlateau(monitor='val_loss', factor=0.5, patience=3,
    ↳ min_lr=1e-6)
```

```
[8]: model.fit(X_train, y_train, epochs=20, validation_data=(X_test, y_test),  
             ↪ batch_size=32, callbacks=[reduce_lr])
```

```
Epoch 1/20  
250/250          1846s 7s/step -  
accuracy: 0.1107 - loss: 3.2723 - val_accuracy: 0.0975 - val_loss: 2.3432 -  
learning_rate: 1.0000e-04  
Epoch 2/20  
250/250          1956s 8s/step -  
accuracy: 0.0971 - loss: 2.3040 - val_accuracy: 0.1055 - val_loss: 2.3580 -  
learning_rate: 1.0000e-04  
Epoch 3/20  
250/250          1642s 7s/step -  
accuracy: 0.1029 - loss: 2.3108 - val_accuracy: 0.0920 - val_loss: 2.3029 -  
learning_rate: 1.0000e-04  
Epoch 4/20  
250/250          1378s 6s/step -  
accuracy: 0.1089 - loss: 2.3031 - val_accuracy: 0.0930 - val_loss: 2.3029 -  
learning_rate: 1.0000e-04  
Epoch 5/20  
250/250          1347s 5s/step -  
accuracy: 0.1057 - loss: 2.3047 - val_accuracy: 0.0930 - val_loss: 2.3030 -  
learning_rate: 1.0000e-04  
Epoch 6/20  
250/250          1301s 5s/step -  
accuracy: 0.1010 - loss: 2.3026 - val_accuracy: 0.0920 - val_loss: 2.3033 -  
learning_rate: 1.0000e-04  
Epoch 7/20  
250/250          1123s 4s/step -  
accuracy: 0.0993 - loss: 2.3022 - val_accuracy: 0.0930 - val_loss: 2.3031 -  
learning_rate: 5.0000e-05  
Epoch 8/20  
250/250          1120s 4s/step -  
accuracy: 0.1059 - loss: 2.3020 - val_accuracy: 0.0915 - val_loss: 2.3031 -  
learning_rate: 5.0000e-05  
Epoch 9/20  
250/250          1138s 4s/step -  
accuracy: 0.0996 - loss: 2.3033 - val_accuracy: 0.0930 - val_loss: 2.3031 -  
learning_rate: 5.0000e-05  
Epoch 10/20  
250/250          1137s 5s/step -  
accuracy: 0.0987 - loss: 2.3031 - val_accuracy: 0.0930 - val_loss: 2.3031 -  
learning_rate: 2.5000e-05  
Epoch 11/20  
250/250          1105s 4s/step -  
accuracy: 0.1025 - loss: 2.3023 - val_accuracy: 0.0925 - val_loss: 2.3031 -  
learning_rate: 2.5000e-05  
Epoch 12/20
```

```

250/250          1099s 4s/step -
accuracy: 0.0992 - loss: 2.3025 - val_accuracy: 0.0930 - val_loss: 2.3031 -
learning_rate: 2.5000e-05
Epoch 13/20
250/250          1127s 4s/step -
accuracy: 0.1038 - loss: 2.3024 - val_accuracy: 0.0930 - val_loss: 2.3031 -
learning_rate: 1.2500e-05
Epoch 14/20
250/250          1022s 4s/step -
accuracy: 0.1034 - loss: 2.3023 - val_accuracy: 0.0975 - val_loss: 2.3026 -
learning_rate: 1.2500e-05
Epoch 15/20
250/250          938s 4s/step -
accuracy: 0.1067 - loss: 2.3024 - val_accuracy: 0.0950 - val_loss: 2.3026 -
learning_rate: 1.2500e-05
Epoch 16/20
250/250          930s 4s/step -
accuracy: 0.1042 - loss: 2.3006 - val_accuracy: 0.1125 - val_loss: 2.2951 -
learning_rate: 1.2500e-05
Epoch 17/20
250/250          904s 4s/step -
accuracy: 0.1118 - loss: 2.2995 - val_accuracy: 0.1530 - val_loss: 2.2610 -
learning_rate: 1.2500e-05
Epoch 18/20
250/250          944s 4s/step -
accuracy: 0.1241 - loss: 2.2888 - val_accuracy: 0.1415 - val_loss: 2.2728 -
learning_rate: 1.2500e-05
Epoch 19/20
250/250          984s 4s/step -
accuracy: 0.1328 - loss: 2.2760 - val_accuracy: 0.1465 - val_loss: 2.2554 -
learning_rate: 1.2500e-05
Epoch 20/20
250/250          927s 4s/step -
accuracy: 0.1410 - loss: 2.2687 - val_accuracy: 0.1705 - val_loss: 2.2497 -
learning_rate: 1.2500e-05

```

[8]: <keras.src.callbacks.history.History at 0x7f074c10af90>

```

[9]: evaluation = model.evaluate(X_test, y_test)
      print(f"Test accuracy: {evaluation[1]:.3f}")

```

```

63/63          29s 466ms/step -
accuracy: 0.1732 - loss: 2.2505
Test accuracy: 0.170

```

4 Apply the confusion matrix after the model


```
[10]: import seaborn as sns
# from sklearn.metrics import confusion
import numpy as NP
from sklearn.metrics import confusion_matrix

cnn_preds = np.argmax(model.predict(X_test), axis=1)
cnn_cm = confusion_matrix(y_test, cnn_preds)

# Plot the confusion matrix
plt.figure(figsize=(8, 6))
sns.heatmap(cnn_cm, annot=True, fmt="d", cmap="Blues", xticklabels=GENRES,
            yticklabels=GENRES)
plt.title("CNN Confusion Matrix")
plt.xlabel("Predicted")
plt.ylabel("True")
plt.show()
```

63/63

28s 443ms/step



4.1 9 - Limited Genres Easy (metal and classical)

```
[11]: import tensorflow as tf
import os
import numpy as np
from sklearn.model_selection import train_test_split

GENRES = ['classical', 'metal']
FILE_PATH = os.path.join('Data', 'tempograms (3 secs)')
X = []
y = []

GENRE_TO_INDEX = {genre: index for index, genre in enumerate(GENRES)}

# Define the augmentation function
def augment_image(image):
    image = tf.image.random_flip_left_right(image)
    image = tf.image.random_brightness(image, max_delta=0.1)
    image = tf.image.random_contrast(image, 0.8, 1.2)
    return image

for genre in GENRES:
    genre_dir = os.path.join(FILE_PATH, genre)
    print(f"Going through {genre}")
    for file in os.listdir(genre_dir):
        image = tf.io.read_file(os.path.join(genre_dir, file))
        image = tf.image.decode_png(image, channels=1)
        image = tf.image.convert_image_dtype(image, tf.float32)
        image = tf.image.resize(image, [256, 256])

        # Apply the augmentation
        image = augment_image(image)
        image = image.numpy() # Convert to numpy array for further processing
        X.append(image)
        y.append(GENRE_TO_INDEX[genre])

X = np.array(X)
y = np.array(y)

# Split the data
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
    ↪random_state=42)

from tensorflow.keras import models
```

```

from tensorflow.keras.layers import Conv2D, MaxPooling2D, Flatten, Dense,
↳Dropout, Normalization

model = models.Sequential([
    Conv2D(32, (3, 3), activation='relu', input_shape=(256, 256, 1)),
    Normalization(),
    MaxPooling2D((2, 2)),

    Conv2D(64, (3, 3), activation='relu'),
    Normalization(),
    MaxPooling2D((2, 2)),

    Conv2D(128, (3, 3), activation='relu'),
    Normalization(),
    MaxPooling2D((2, 2)),

    Conv2D(256, (3, 3), activation='relu'),
    Normalization(),
    MaxPooling2D((2, 2)),

    Flatten(),

    Dense(512, activation='relu'),
    Dropout(0.5),

    Dense(256, activation='relu'),
    Dropout(0.5),

    Dense(128, activation='relu'),
    Dense(10, activation='softmax')
])

from tensorflow.keras.optimizers import Adam
from tensorflow.keras.callbacks import ReduceLROnPlateau

model.compile(optimizer=Adam(learning_rate=0.0001),
↳loss='sparse_categorical_crossentropy', metrics=['accuracy'])
reduce_lr = ReduceLROnPlateau(monitor='val_loss', factor=0.5, patience=3,
↳min_lr=1e-6)

model.fit(X_train, y_train, epochs=20, validation_data=(X_test, y_test),
↳batch_size=32, callbacks=[reduce_lr])

evaluation = model.evaluate(X_test, y_test)
print(f"Test accuracy: {evaluation[1]:.3f}")

```

Going through classical

Going through metal

```
/opt/conda/lib/python3.12/site-  
packages/keras/src/layers/convolutional/base_conv.py:107: UserWarning: Do not  
pass an `input_shape`/`input_dim` argument to a layer. When using Sequential  
models, prefer using an `Input(shape)` object as the first layer in the model  
instead.
```

```
    super().__init__(activity_regularizer=activity_regularizer, **kwargs)
```

Epoch 1/20

```
50/50          87s 2s/step -  
accuracy: 0.4488 - loss: 1.6158 - val_accuracy: 0.4875 - val_loss: 0.7366 -  
learning_rate: 1.0000e-04
```

Epoch 2/20

```
50/50          86s 2s/step -  
accuracy: 0.4632 - loss: 0.8819 - val_accuracy: 0.5025 - val_loss: 0.7149 -  
learning_rate: 1.0000e-04
```

Epoch 3/20

```
50/50          92s 2s/step -  
accuracy: 0.5114 - loss: 0.7963 - val_accuracy: 0.5025 - val_loss: 0.7058 -  
learning_rate: 1.0000e-04
```

Epoch 4/20

```
50/50         151s 2s/step -  
accuracy: 0.5119 - loss: 0.7515 - val_accuracy: 0.5025 - val_loss: 0.7045 -  
learning_rate: 1.0000e-04
```

Epoch 5/20

```
50/50          94s 2s/step -  
accuracy: 0.4782 - loss: 0.7688 - val_accuracy: 0.4975 - val_loss: 0.6977 -  
learning_rate: 1.0000e-04
```

Epoch 6/20

```
50/50         161s 2s/step -  
accuracy: 0.4965 - loss: 0.7432 - val_accuracy: 0.4975 - val_loss: 0.6988 -  
learning_rate: 1.0000e-04
```

Epoch 7/20

```
50/50          97s 2s/step -  
accuracy: 0.5427 - loss: 0.7220 - val_accuracy: 0.4975 - val_loss: 0.6972 -  
learning_rate: 1.0000e-04
```

Epoch 8/20

```
50/50         107s 2s/step -  
accuracy: 0.5343 - loss: 0.7180 - val_accuracy: 0.4975 - val_loss: 0.6979 -  
learning_rate: 1.0000e-04
```

Epoch 9/20

```
50/50         108s 2s/step -  
accuracy: 0.5194 - loss: 0.7140 - val_accuracy: 0.5025 - val_loss: 0.6963 -  
learning_rate: 1.0000e-04
```

Epoch 10/20

```
50/50         107s 2s/step -  
accuracy: 0.4949 - loss: 0.7308 - val_accuracy: 0.4950 - val_loss: 0.6948 -  
learning_rate: 1.0000e-04
```

Epoch 11/20
50/50 104s 2s/step -
accuracy: 0.5241 - loss: 0.7153 - val_accuracy: 0.5025 - val_loss: 0.7000 -
learning_rate: 1.0000e-04

Epoch 12/20
50/50 106s 2s/step -
accuracy: 0.4997 - loss: 0.7124 - val_accuracy: 0.5425 - val_loss: 0.6928 -
learning_rate: 1.0000e-04

Epoch 13/20
50/50 129s 2s/step -
accuracy: 0.4742 - loss: 0.7269 - val_accuracy: 0.5025 - val_loss: 0.6970 -
learning_rate: 1.0000e-04

Epoch 14/20
50/50 94s 2s/step -
accuracy: 0.5180 - loss: 0.7123 - val_accuracy: 0.4975 - val_loss: 0.6932 -
learning_rate: 1.0000e-04

Epoch 15/20
50/50 100s 2s/step -
accuracy: 0.4656 - loss: 0.7232 - val_accuracy: 0.5175 - val_loss: 0.6935 -
learning_rate: 1.0000e-04

Epoch 16/20
50/50 136s 2s/step -
accuracy: 0.5044 - loss: 0.7168 - val_accuracy: 0.5025 - val_loss: 0.6926 -
learning_rate: 5.0000e-05

Epoch 17/20
50/50 129s 2s/step -
accuracy: 0.5048 - loss: 0.7062 - val_accuracy: 0.6400 - val_loss: 0.6936 -
learning_rate: 5.0000e-05

Epoch 18/20
50/50 60s 1s/step -
accuracy: 0.5059 - loss: 0.7114 - val_accuracy: 0.6400 - val_loss: 0.6929 -
learning_rate: 5.0000e-05

Epoch 19/20
50/50 85s 2s/step -
accuracy: 0.5209 - loss: 0.7084 - val_accuracy: 0.5375 - val_loss: 0.6937 -
learning_rate: 5.0000e-05

Epoch 20/20
50/50 87s 2s/step -
accuracy: 0.4852 - loss: 0.7141 - val_accuracy: 0.5375 - val_loss: 0.6924 -
learning_rate: 2.5000e-05

13/13 5s 381ms/step -
accuracy: 0.5281 - loss: 0.6925

Test accuracy: 0.538

4.2 10 - Confusion Matrix Easy (classical and metal)


```
[12]: import seaborn as sns
# from sklearn.metrics import confusion
import numpy as NP
from sklearn.metrics import confusion_matrix

cnn_preds = np.argmax(model.predict(X_test), axis=1)
cnn_cm = confusion_matrix(y_test, cnn_preds)

# Plot the confusion matrix
plt.figure(figsize=(8, 6))
sns.heatmap(cnn_cm, annot=True, fmt="d", cmap="Blues", xticklabels=GENRES, yticklabels=GENRES)
plt.title("CNN Confusion Matrix")
plt.xlabel("Predicted")
plt.ylabel("True")
plt.show()
```

13/13

5s 391ms/step



4.3 11 - Limited genres Hard (disco and pop)

```
[13]: import tensorflow as tf
import os
import numpy as np
from sklearn.model_selection import train_test_split

GENRES = ['disco', 'pop']
FILE_PATH = os.path.join('Data', 'tempograms (3 secs)')
X = []
y = []

GENRE_TO_INDEX = {genre: index for index, genre in enumerate(GENRES)}

# Define the augmentation function
def augment_image(image):
    image = tf.image.random_flip_left_right(image)
    image = tf.image.random_brightness(image, max_delta=0.1)
    image = tf.image.random_contrast(image, 0.8, 1.2)
    return image

for genre in GENRES:
    genre_dir = os.path.join(FILE_PATH, genre)
    print(f"Going through {genre}")
    for file in os.listdir(genre_dir):
        image = tf.io.read_file(os.path.join(genre_dir, file))
        image = tf.image.decode_png(image, channels=1)
        image = tf.image.convert_image_dtype(image, tf.float32)
        image = tf.image.resize(image, [256, 256])

        # Apply the augmentation
        image = augment_image(image)

        image = image.numpy() # Convert to numpy array for further processing
        X.append(image)
        y.append(GENRE_TO_INDEX[genre])

X = np.array(X)
y = np.array(y)

# Split the data
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
    ↪random_state=42)
```

```

from tensorflow.keras import models
from tensorflow.keras.layers import Conv2D, MaxPooling2D, Flatten, Dense,
↳Dropout, Normalization

model = models.Sequential([
    Conv2D(32, (3, 3), activation='relu', input_shape=(256, 256, 1)),
    Normalization(),
    MaxPooling2D((2, 2)),

    Conv2D(64, (3, 3), activation='relu'),
    Normalization(),
    MaxPooling2D((2, 2)),

    Conv2D(128, (3, 3), activation='relu'),
    Normalization(),
    MaxPooling2D((2, 2)),

    Conv2D(256, (3, 3), activation='relu'),
    Normalization(),
    MaxPooling2D((2, 2)),

    Flatten(),

    Dense(512, activation='relu'),
    Dropout(0.5),

    Dense(256, activation='relu'),
    Dropout(0.5),

    Dense(128, activation='relu'),
    Dense(10, activation='softmax')
])

from tensorflow.keras.optimizers import Adam
from tensorflow.keras.callbacks import ReduceLROnPlateau

model.compile(optimizer=Adam(learning_rate=0.0001),
↳loss='sparse_categorical_crossentropy', metrics=['accuracy'])
reduce_lr = ReduceLROnPlateau(monitor='val_loss', factor=0.5, patience=3,
↳min_lr=1e-6)

model.fit(X_train, y_train, epochs=20, validation_data=(X_test, y_test),
↳batch_size=32, callbacks=[reduce_lr])

evaluation = model.evaluate(X_test, y_test)
print(f"Test accuracy: {evaluation[1]:.3f}")

```

Going through disco

Going through pop

/opt/conda/lib/python3.12/site-

packages/keras/src/layers/convolutional/base_conv.py:107: UserWarning: Do not pass an `input_shape`/`input_dim` argument to a layer. When using Sequential models, prefer using an `Input(shape)` object as the first layer in the model instead.

```
super().__init__(activity_regularizer=activity_regularizer, **kwargs)
```

Epoch 1/20

50/50 93s 2s/step -

accuracy: 0.4358 - loss: 1.5702 - val_accuracy: 0.5025 - val_loss: 0.7514 -
learning_rate: 1.0000e-04

Epoch 2/20

50/50 85s 2s/step -

accuracy: 0.5055 - loss: 0.8618 - val_accuracy: 0.4975 - val_loss: 0.7315 -
learning_rate: 1.0000e-04

Epoch 3/20

50/50 79s 2s/step -

accuracy: 0.5094 - loss: 0.8087 - val_accuracy: 0.6000 - val_loss: 0.7177 -
learning_rate: 1.0000e-04

Epoch 4/20

50/50 68s 1s/step -

accuracy: 0.5013 - loss: 0.7872 - val_accuracy: 0.5025 - val_loss: 0.7353 -
learning_rate: 1.0000e-04

Epoch 5/20

50/50 76s 2s/step -

accuracy: 0.5387 - loss: 0.7382 - val_accuracy: 0.5025 - val_loss: 0.7043 -
learning_rate: 1.0000e-04

Epoch 6/20

50/50 85s 2s/step -

accuracy: 0.4790 - loss: 0.7619 - val_accuracy: 0.4975 - val_loss: 0.7018 -
learning_rate: 1.0000e-04

Epoch 7/20

50/50 77s 2s/step -

accuracy: 0.4961 - loss: 0.7465 - val_accuracy: 0.4975 - val_loss: 0.7105 -
learning_rate: 1.0000e-04

Epoch 8/20

50/50 78s 2s/step -

accuracy: 0.4984 - loss: 0.7297 - val_accuracy: 0.4975 - val_loss: 0.6993 -
learning_rate: 1.0000e-04

Epoch 9/20

50/50 82s 2s/step -

accuracy: 0.5192 - loss: 0.7251 - val_accuracy: 0.5900 - val_loss: 0.6970 -
learning_rate: 1.0000e-04

Epoch 10/20

50/50 74s 1s/step -

accuracy: 0.5024 - loss: 0.7182 - val_accuracy: 0.4975 - val_loss: 0.6926 -


```

learning_rate: 1.0000e-04
Epoch 11/20
50/50          86s 2s/step -
accuracy: 0.5147 - loss: 0.7292 - val_accuracy: 0.5825 - val_loss: 0.6795 -
learning_rate: 1.0000e-04
Epoch 12/20
50/50          86s 2s/step -
accuracy: 0.5181 - loss: 0.7168 - val_accuracy: 0.6025 - val_loss: 0.6741 -
learning_rate: 1.0000e-04
Epoch 13/20
50/50          85s 2s/step -
accuracy: 0.5525 - loss: 0.7013 - val_accuracy: 0.6425 - val_loss: 0.6603 -
learning_rate: 1.0000e-04
Epoch 14/20
50/50          78s 2s/step -
accuracy: 0.5737 - loss: 0.6865 - val_accuracy: 0.6050 - val_loss: 0.6577 -
learning_rate: 1.0000e-04
Epoch 15/20
50/50          77s 2s/step -
accuracy: 0.6078 - loss: 0.6641 - val_accuracy: 0.6725 - val_loss: 0.6176 -
learning_rate: 1.0000e-04
Epoch 16/20
50/50          76s 2s/step -
accuracy: 0.6471 - loss: 0.6437 - val_accuracy: 0.7100 - val_loss: 0.6101 -
learning_rate: 1.0000e-04
Epoch 17/20
50/50          80s 2s/step -
accuracy: 0.6867 - loss: 0.6134 - val_accuracy: 0.7350 - val_loss: 0.5834 -
learning_rate: 1.0000e-04
Epoch 18/20
50/50          73s 1s/step -
accuracy: 0.7021 - loss: 0.5804 - val_accuracy: 0.6925 - val_loss: 0.6476 -
learning_rate: 1.0000e-04
Epoch 19/20
50/50          82s 2s/step -
accuracy: 0.7163 - loss: 0.5865 - val_accuracy: 0.7525 - val_loss: 0.5558 -
learning_rate: 1.0000e-04
Epoch 20/20
50/50          74s 1s/step -
accuracy: 0.7337 - loss: 0.5382 - val_accuracy: 0.7300 - val_loss: 0.5693 -
learning_rate: 1.0000e-04
13/13          4s 283ms/step -
accuracy: 0.7387 - loss: 0.5738
Test accuracy: 0.730

```

4.4 12 - Confusion Matrix Hard (disco and pop)


```
[14]: import seaborn as sns
# from sklearn.metrics import confusion
import numpy as NP
from sklearn.metrics import confusion_matrix

cnn_preds = np.argmax(model.predict(X_test), axis=1)
cnn_cm = confusion_matrix(y_test, cnn_preds)

# Plot the confusion matrix
plt.figure(figsize=(8, 6))
sns.heatmap(cnn_cm, annot=True, fmt="d", cmap="Blues", xticklabels=GENRES, yticklabels=GENRES)
plt.title("CNN Confusion Matrix")
plt.xlabel("Predicted")
plt.ylabel("True")
plt.show()
```

13/13

5s 379ms/step



4.5 13 - Limited Genres Medium (5 random)

```
[15]: import tensorflow as tf
import os
import numpy as np
from sklearn.model_selection import train_test_split
import random

GENRES = ['blues', 'classical', 'country', 'disco', 'hiphop', 'jazz', 'metal', 'pop', 'reggae', 'rock']
GENRES = random.sample(GENRES, 5)
print(GENRES)
FILE_PATH = os.path.join('Data', 'tempograms (3 secs)')
X = []
y = []

GENRE_TO_INDEX = {genre: index for index, genre in enumerate(GENRES)}

# Define the augmentation function
def augment_image(image):
    image = tf.image.random_flip_left_right(image)
    image = tf.image.random_brightness(image, max_delta=0.1)
    image = tf.image.random_contrast(image, 0.8, 1.2)
    return image

for genre in GENRES:
    genre_dir = os.path.join(FILE_PATH, genre)
    print(f"Going through {genre}")
    for file in os.listdir(genre_dir):
        image = tf.io.read_file(os.path.join(genre_dir, file))
        image = tf.image.decode_png(image, channels=1)
        image = tf.image.convert_image_dtype(image, tf.float32)
        image = tf.image.resize(image, [256, 256])

        # Apply the augmentation
        image = augment_image(image)

        image = image.numpy() # Convert to numpy array for further processing
        X.append(image)
        y.append(GENRE_TO_INDEX[genre])

X = np.array(X)
y = np.array(y)
```

```

# Split the data
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
    ↪random_state=42)

from tensorflow.keras import models
from tensorflow.keras.layers import Conv2D, MaxPooling2D, Flatten, Dense,
    ↪Dropout, Normalization

model = models.Sequential([
    Conv2D(32, (3, 3), activation='relu', input_shape=(256, 256, 1)),
    Normalization(),
    MaxPooling2D((2, 2)),

    Conv2D(64, (3, 3), activation='relu'),
    Normalization(),
    MaxPooling2D((2, 2)),

    Conv2D(128, (3, 3), activation='relu'),
    Normalization(),
    MaxPooling2D((2, 2)),

    Conv2D(256, (3, 3), activation='relu'),
    Normalization(),
    MaxPooling2D((2, 2)),

    Flatten(),

    Dense(512, activation='relu'),
    Dropout(0.5),

    Dense(256, activation='relu'),
    Dropout(0.5),

    Dense(128, activation='relu'),
    Dense(10, activation='softmax')
])

from tensorflow.keras.optimizers import Adam
from tensorflow.keras.callbacks import ReduceLROnPlateau

model.compile(optimizer=Adam(learning_rate=0.0001),
    ↪loss='sparse_categorical_crossentropy', metrics=['accuracy'])
reduce_lr = ReduceLROnPlateau(monitor='val_loss', factor=0.5, patience=3,
    ↪min_lr=1e-6)

model.fit(X_train, y_train, epochs=20, validation_data=(X_test, y_test),
    ↪batch_size=32, callbacks=[reduce_lr])

```

```
evaluation = model.evaluate(X_test, y_test)
print(f"Test accuracy: {evaluation[1]:.3f}")
```

```
['metal', 'disco', 'jazz', 'classical', 'rock']
```

```
Going through metal
```

```
Going through disco
```

```
Going through jazz
```

```
Going through classical
```

```
Going through rock
```

```
/opt/conda/lib/python3.12/site-
```

```
packages/keras/src/layers/convolutional/base_conv.py:107: UserWarning: Do not
pass an `input_shape`/`input_dim` argument to a layer. When using Sequential
models, prefer using an `Input(shape)` object as the first layer in the model
instead.
```

```
super().__init__(activity_regularizer=activity_regularizer, **kwargs)
```

```
Epoch 1/20
```

```
125/125          215s 2s/step -
```

```
accuracy: 0.1996 - loss: 1.9834 - val_accuracy: 0.1920 - val_loss: 1.6781 -
```

```
learning_rate: 1.0000e-04
```

```
Epoch 2/20
```

```
125/125          229s 1s/step -
```

```
accuracy: 0.2443 - loss: 1.6729 - val_accuracy: 0.3200 - val_loss: 1.5157 -
```

```
learning_rate: 1.0000e-04
```

```
Epoch 3/20
```

```
125/125          170s 1s/step -
```

```
accuracy: 0.3162 - loss: 1.5440 - val_accuracy: 0.4110 - val_loss: 1.4048 -
```

```
learning_rate: 1.0000e-04
```

```
Epoch 4/20
```

```
125/125          208s 2s/step -
```

```
accuracy: 0.3493 - loss: 1.4838 - val_accuracy: 0.4010 - val_loss: 1.4004 -
```

```
learning_rate: 1.0000e-04
```

```
Epoch 5/20
```

```
125/125          212s 2s/step -
```

```
accuracy: 0.3532 - loss: 1.4689 - val_accuracy: 0.4090 - val_loss: 1.3775 -
```

```
learning_rate: 1.0000e-04
```

```
Epoch 6/20
```

```
125/125          202s 2s/step -
```

```
accuracy: 0.3750 - loss: 1.4490 - val_accuracy: 0.4240 - val_loss: 1.3526 -
```

```
learning_rate: 1.0000e-04
```

```
Epoch 7/20
```

```
125/125          213s 2s/step -
```

```
accuracy: 0.3818 - loss: 1.4062 - val_accuracy: 0.4440 - val_loss: 1.3735 -
```

```
learning_rate: 1.0000e-04
```

```
Epoch 8/20
```

```
125/125          214s 2s/step -
```

```
accuracy: 0.3957 - loss: 1.3938 - val_accuracy: 0.4540 - val_loss: 1.3376 -
```

```

learning_rate: 1.0000e-04
Epoch 9/20
125/125          203s 2s/step -
accuracy: 0.4173 - loss: 1.3556 - val_accuracy: 0.4480 - val_loss: 1.3074 -
learning_rate: 1.0000e-04
Epoch 10/20
125/125          195s 2s/step -
accuracy: 0.4193 - loss: 1.3393 - val_accuracy: 0.4600 - val_loss: 1.2939 -
learning_rate: 1.0000e-04
Epoch 11/20
125/125          210s 2s/step -
accuracy: 0.4387 - loss: 1.3117 - val_accuracy: 0.4470 - val_loss: 1.3248 -
learning_rate: 1.0000e-04
Epoch 12/20
125/125          203s 2s/step -
accuracy: 0.4431 - loss: 1.3166 - val_accuracy: 0.4390 - val_loss: 1.3498 -
learning_rate: 1.0000e-04
Epoch 13/20
125/125          214s 2s/step -
accuracy: 0.4466 - loss: 1.3065 - val_accuracy: 0.4780 - val_loss: 1.2788 -
learning_rate: 1.0000e-04
Epoch 14/20
125/125          211s 2s/step -
accuracy: 0.4672 - loss: 1.2686 - val_accuracy: 0.4790 - val_loss: 1.2579 -
learning_rate: 1.0000e-04
Epoch 15/20
125/125          213s 2s/step -
accuracy: 0.4530 - loss: 1.2769 - val_accuracy: 0.4520 - val_loss: 1.3249 -
learning_rate: 1.0000e-04
Epoch 16/20
125/125          214s 2s/step -
accuracy: 0.4750 - loss: 1.2655 - val_accuracy: 0.4910 - val_loss: 1.2537 -
learning_rate: 1.0000e-04
Epoch 17/20
125/125          200s 2s/step -
accuracy: 0.4596 - loss: 1.2684 - val_accuracy: 0.4950 - val_loss: 1.2662 -
learning_rate: 1.0000e-04
Epoch 18/20
125/125          213s 2s/step -
accuracy: 0.4943 - loss: 1.2298 - val_accuracy: 0.5020 - val_loss: 1.2408 -
learning_rate: 1.0000e-04
Epoch 19/20
125/125          259s 2s/step -
accuracy: 0.4795 - loss: 1.2475 - val_accuracy: 0.5160 - val_loss: 1.2006 -
learning_rate: 1.0000e-04
Epoch 20/20
125/125          201s 2s/step -
accuracy: 0.4924 - loss: 1.2366 - val_accuracy: 0.5120 - val_loss: 1.2230 -

```

```
learning_rate: 1.0000e-04
32/32          12s 372ms/step -
accuracy: 0.5205 - loss: 1.1956
Test accuracy: 0.512
```

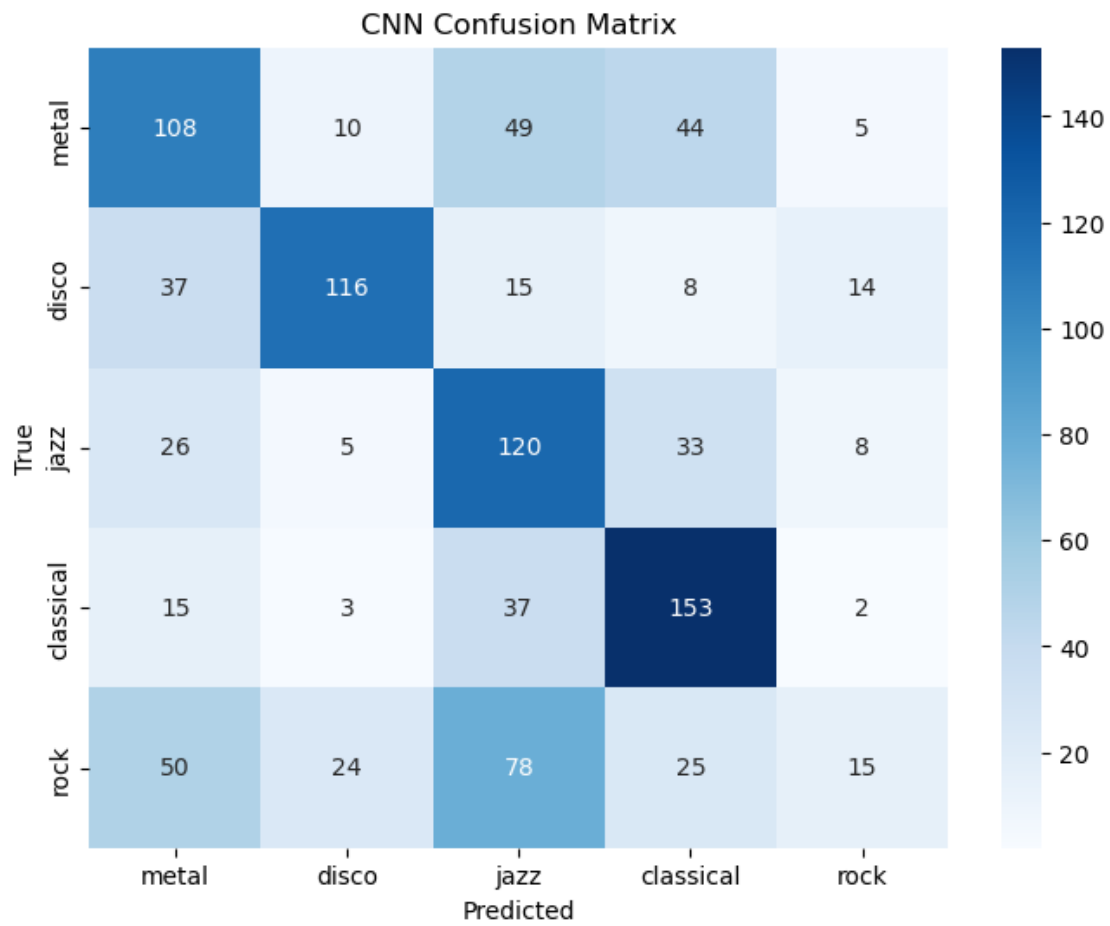
4.6 14 - Confusion Matrix Medium (5 random)

```
[16]: import seaborn as sns
      # from sklearn.metrics import confusion
      import numpy as NP
      from sklearn.metrics import confusion_matrix

      cnn_preds = np.argmax(model.predict(X_test), axis=1)
      cnn_cm = confusion_matrix(y_test, cnn_preds)

      # Plot the confusion matrix
      plt.figure(figsize=(8, 6))
      sns.heatmap(cnn_cm, annot=True, fmt="d", cmap="Blues", xticklabels=GENRES,
        ↪yticklabels=GENRES)
      plt.title("CNN Confusion Matrix")
      plt.xlabel("Predicted")
      plt.ylabel("True")
      plt.show()
```

```
32/32          13s 390ms/step
```



[]: